

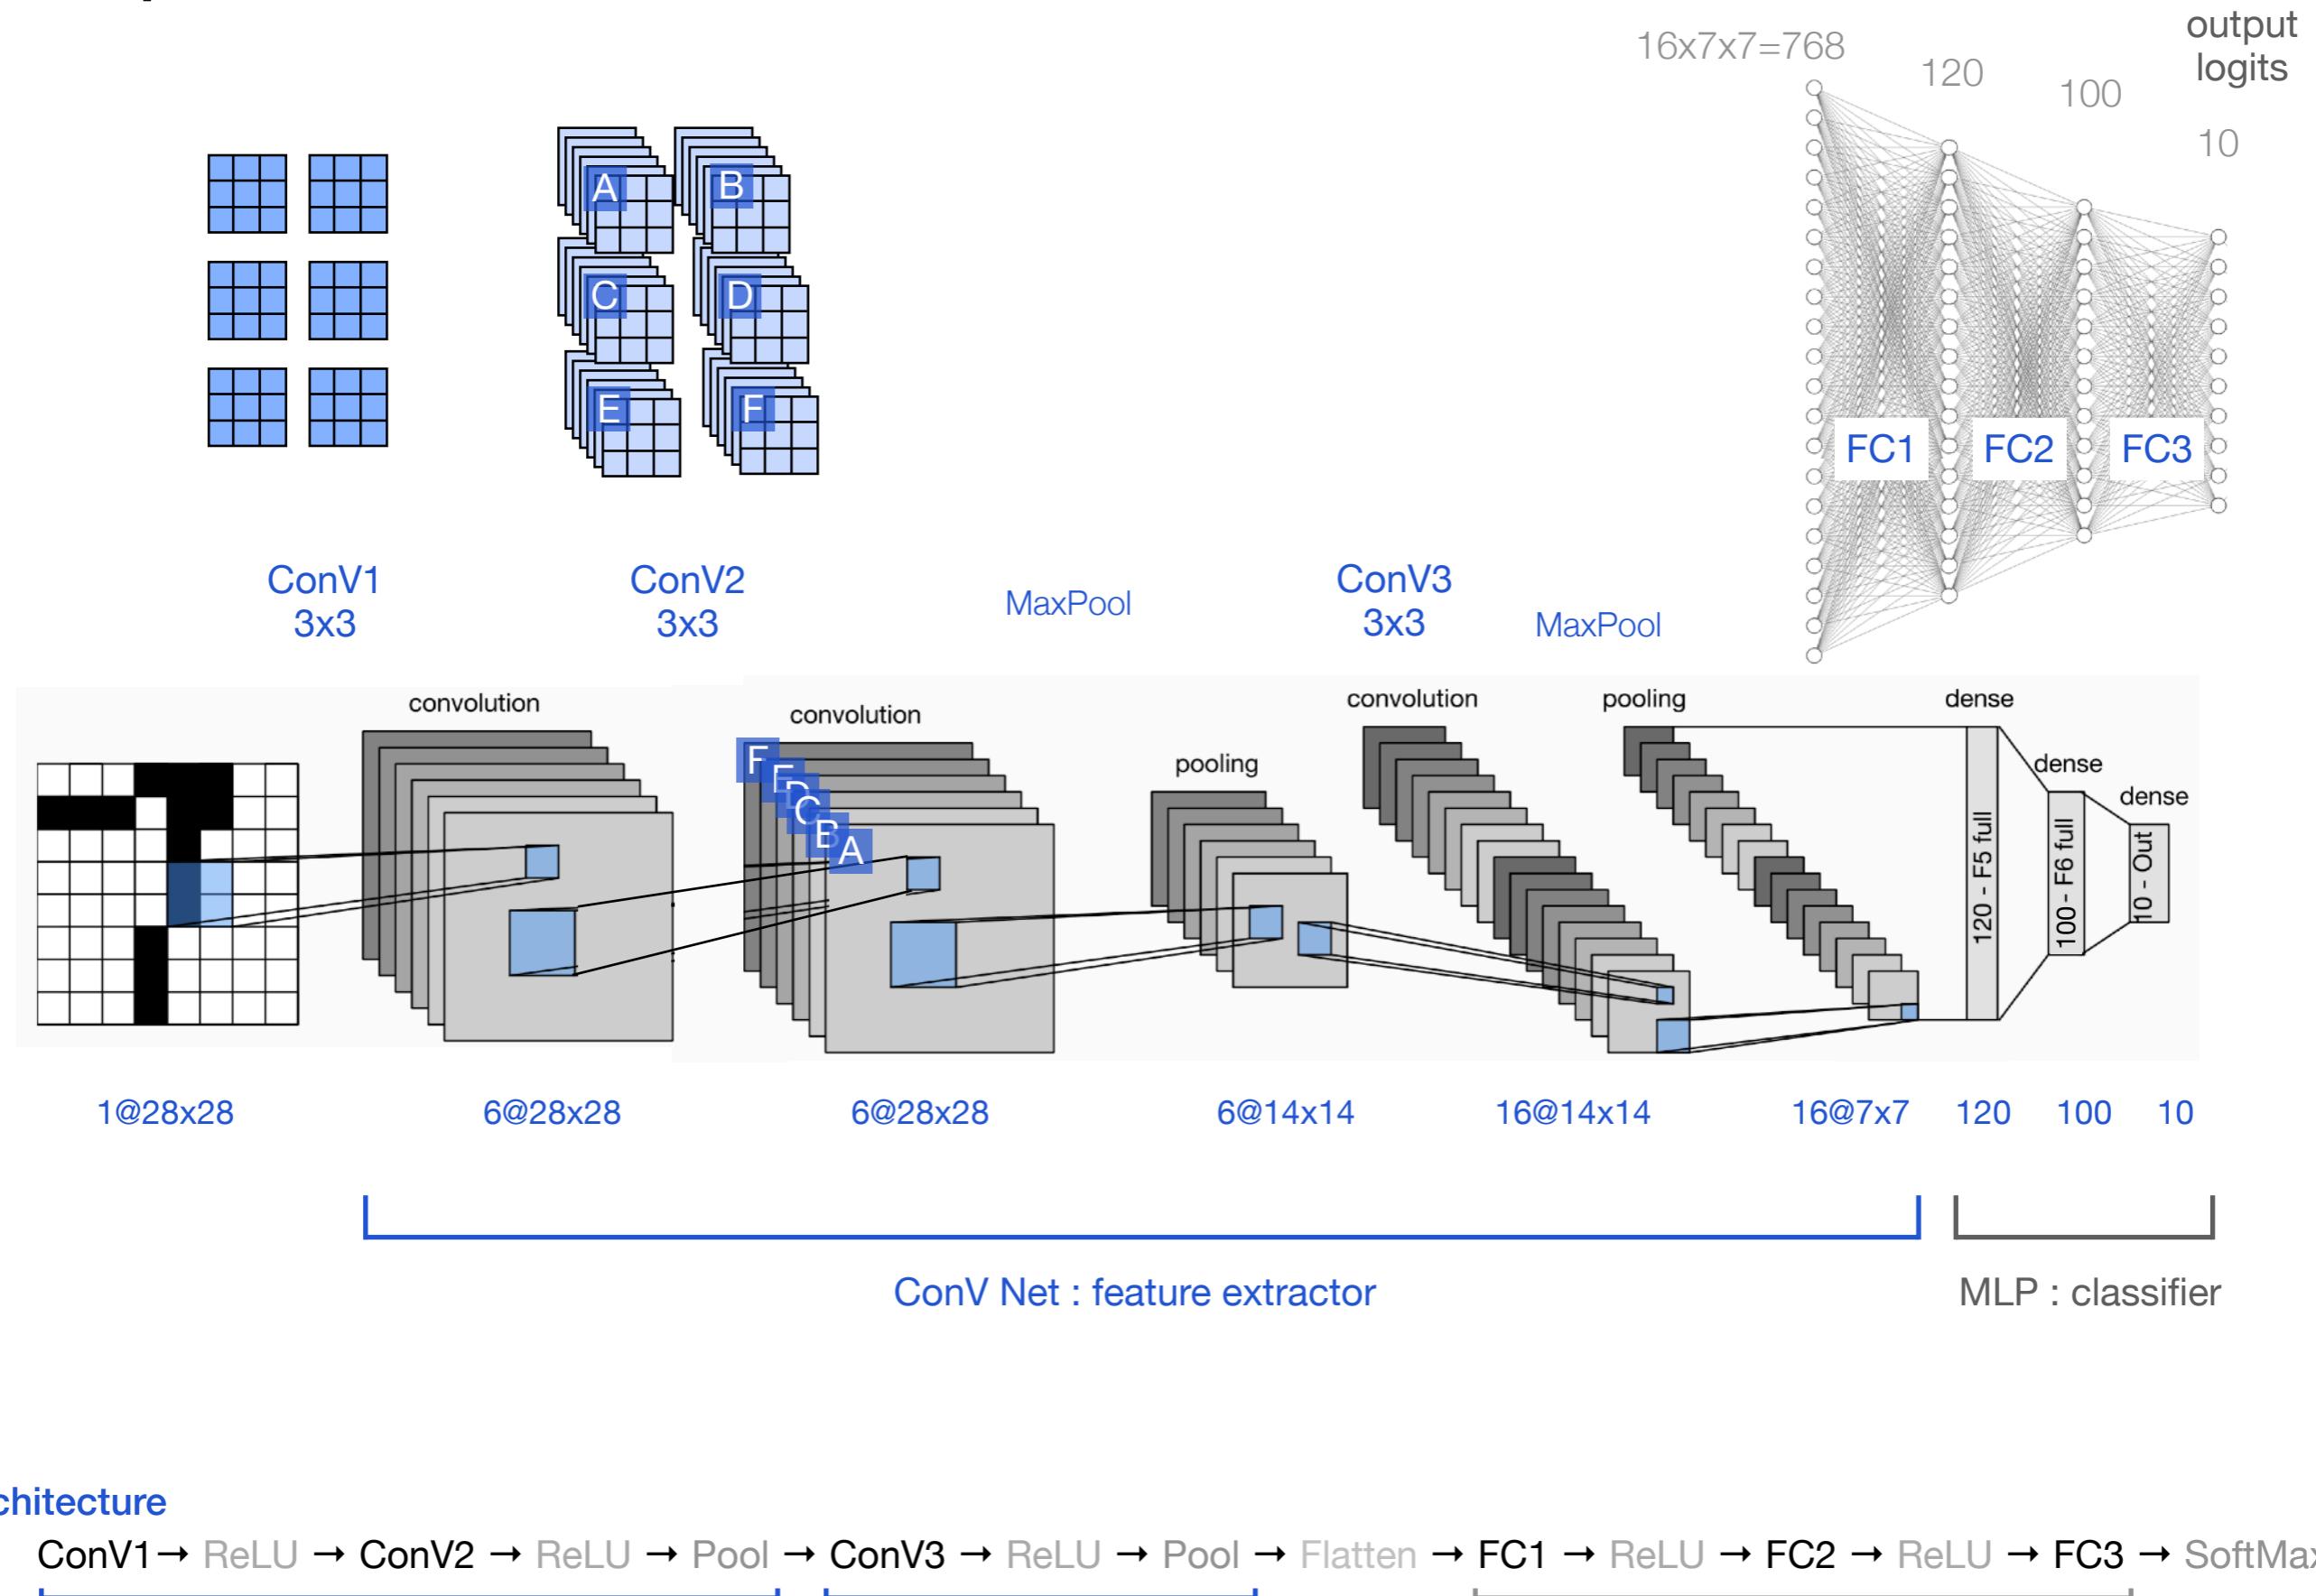
Exploring Features Learned from Convolutional Networks

Outline

- A simple CNN trained on fashion-MNIST dataset
- Convolutional filters learned from CNNs
- Feature Vectors learned in the Last Layer | UMAP visualization
- Transfer Learning



A simple ConV Net



Fashion-MNIST

0 : T-shirt

1 : Trouser

2 : Pullover

3 : Dress

4 : Coat

5 : Sandal

6 : Shirt

7 : Sneaker

8 : Bag

9 : Ankle boot

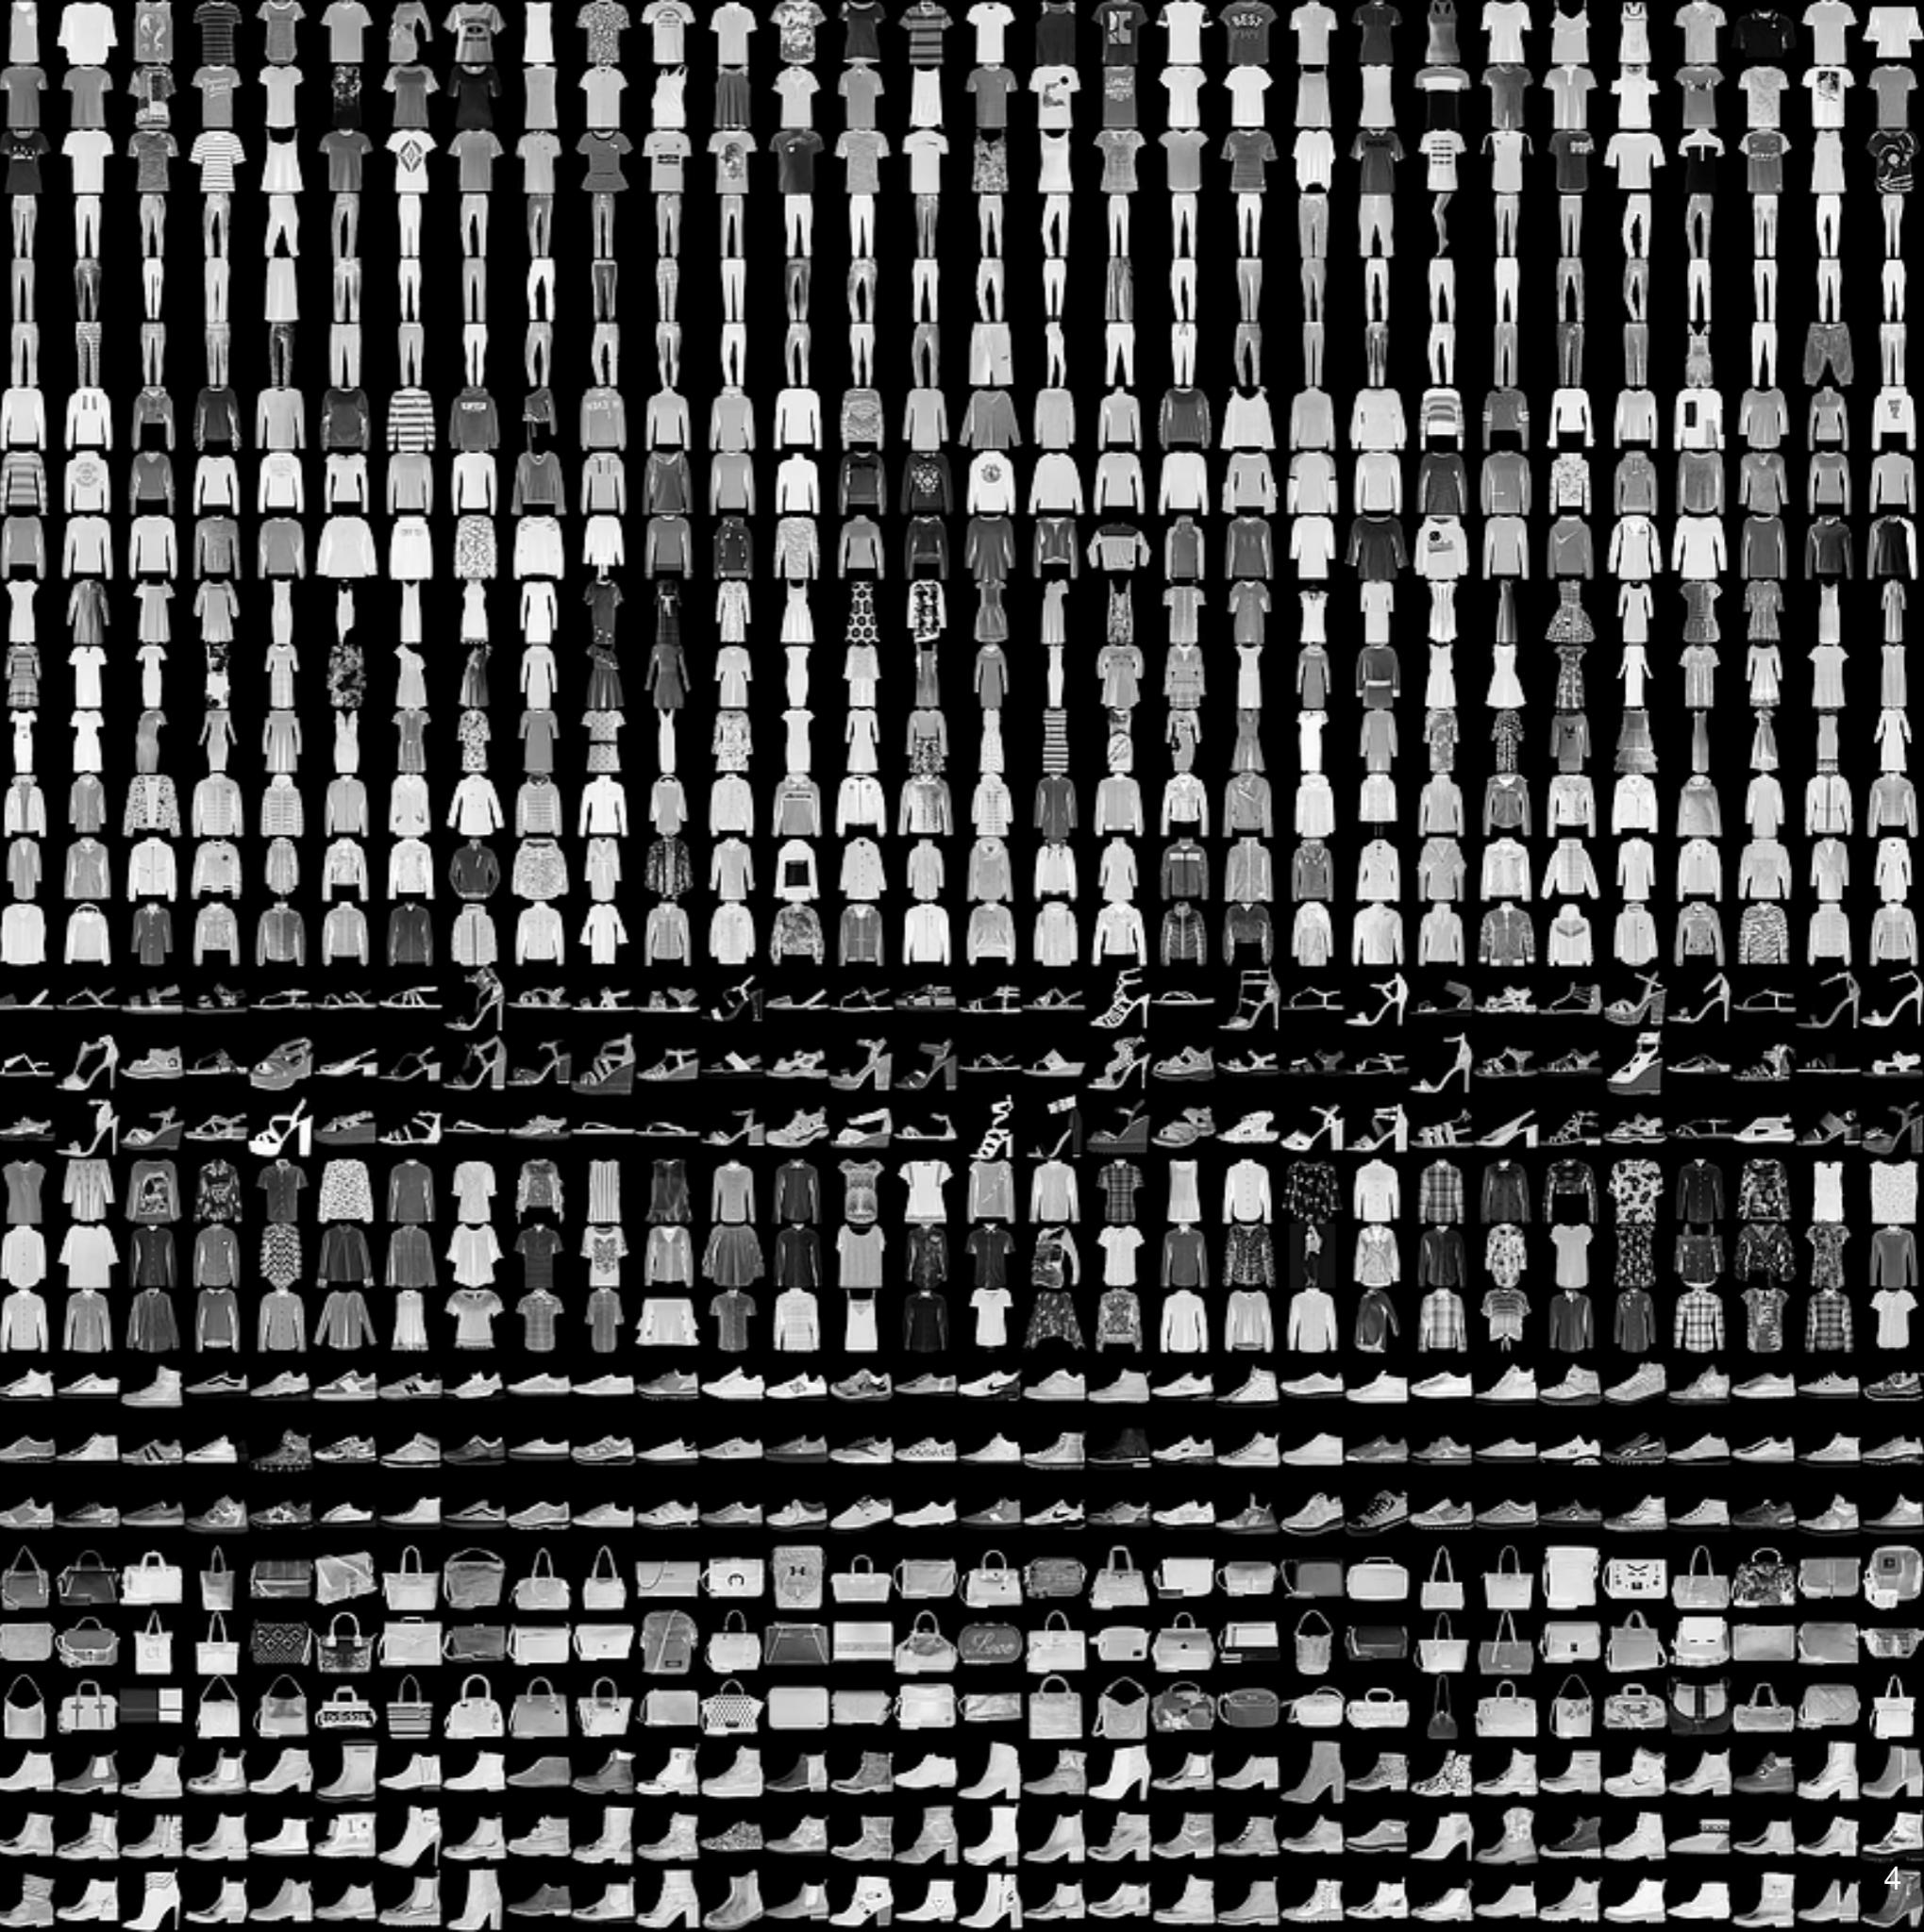


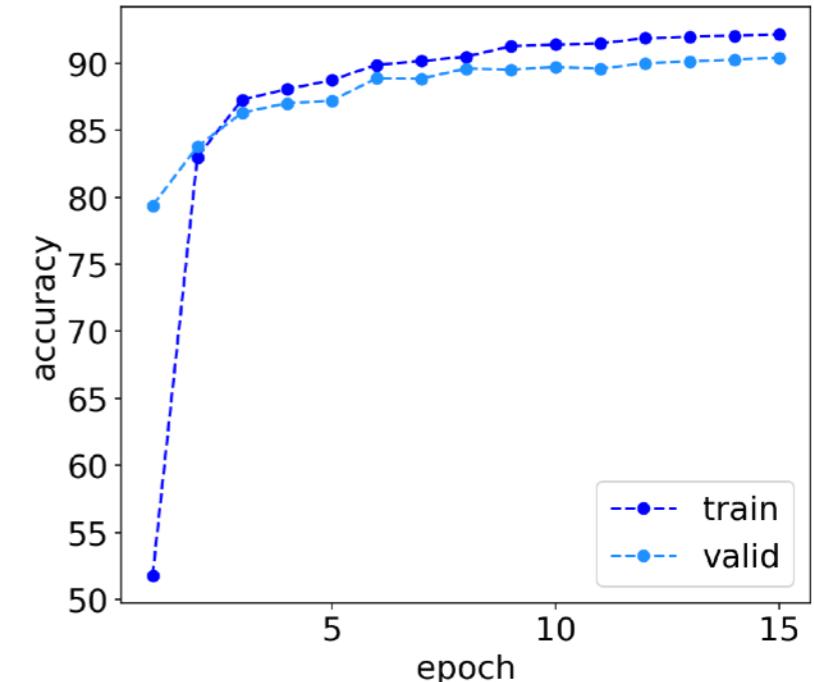
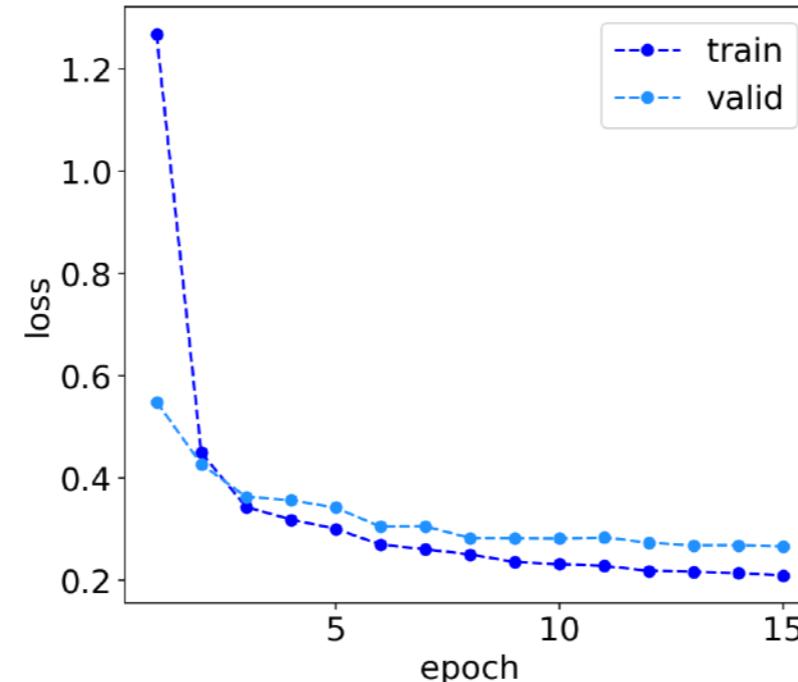
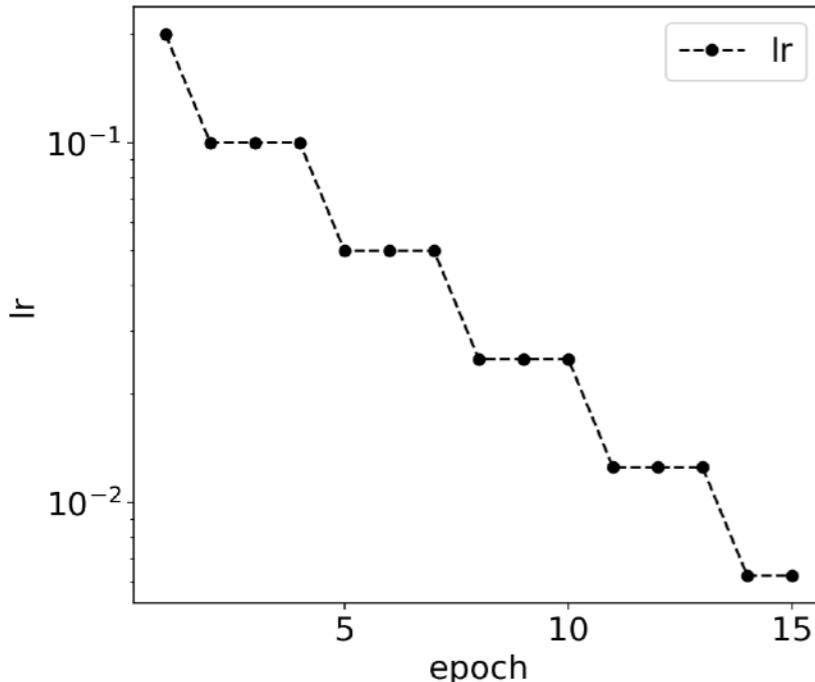
Image size : 28x28

Training Set : 60,000

Testing Set : 10,000

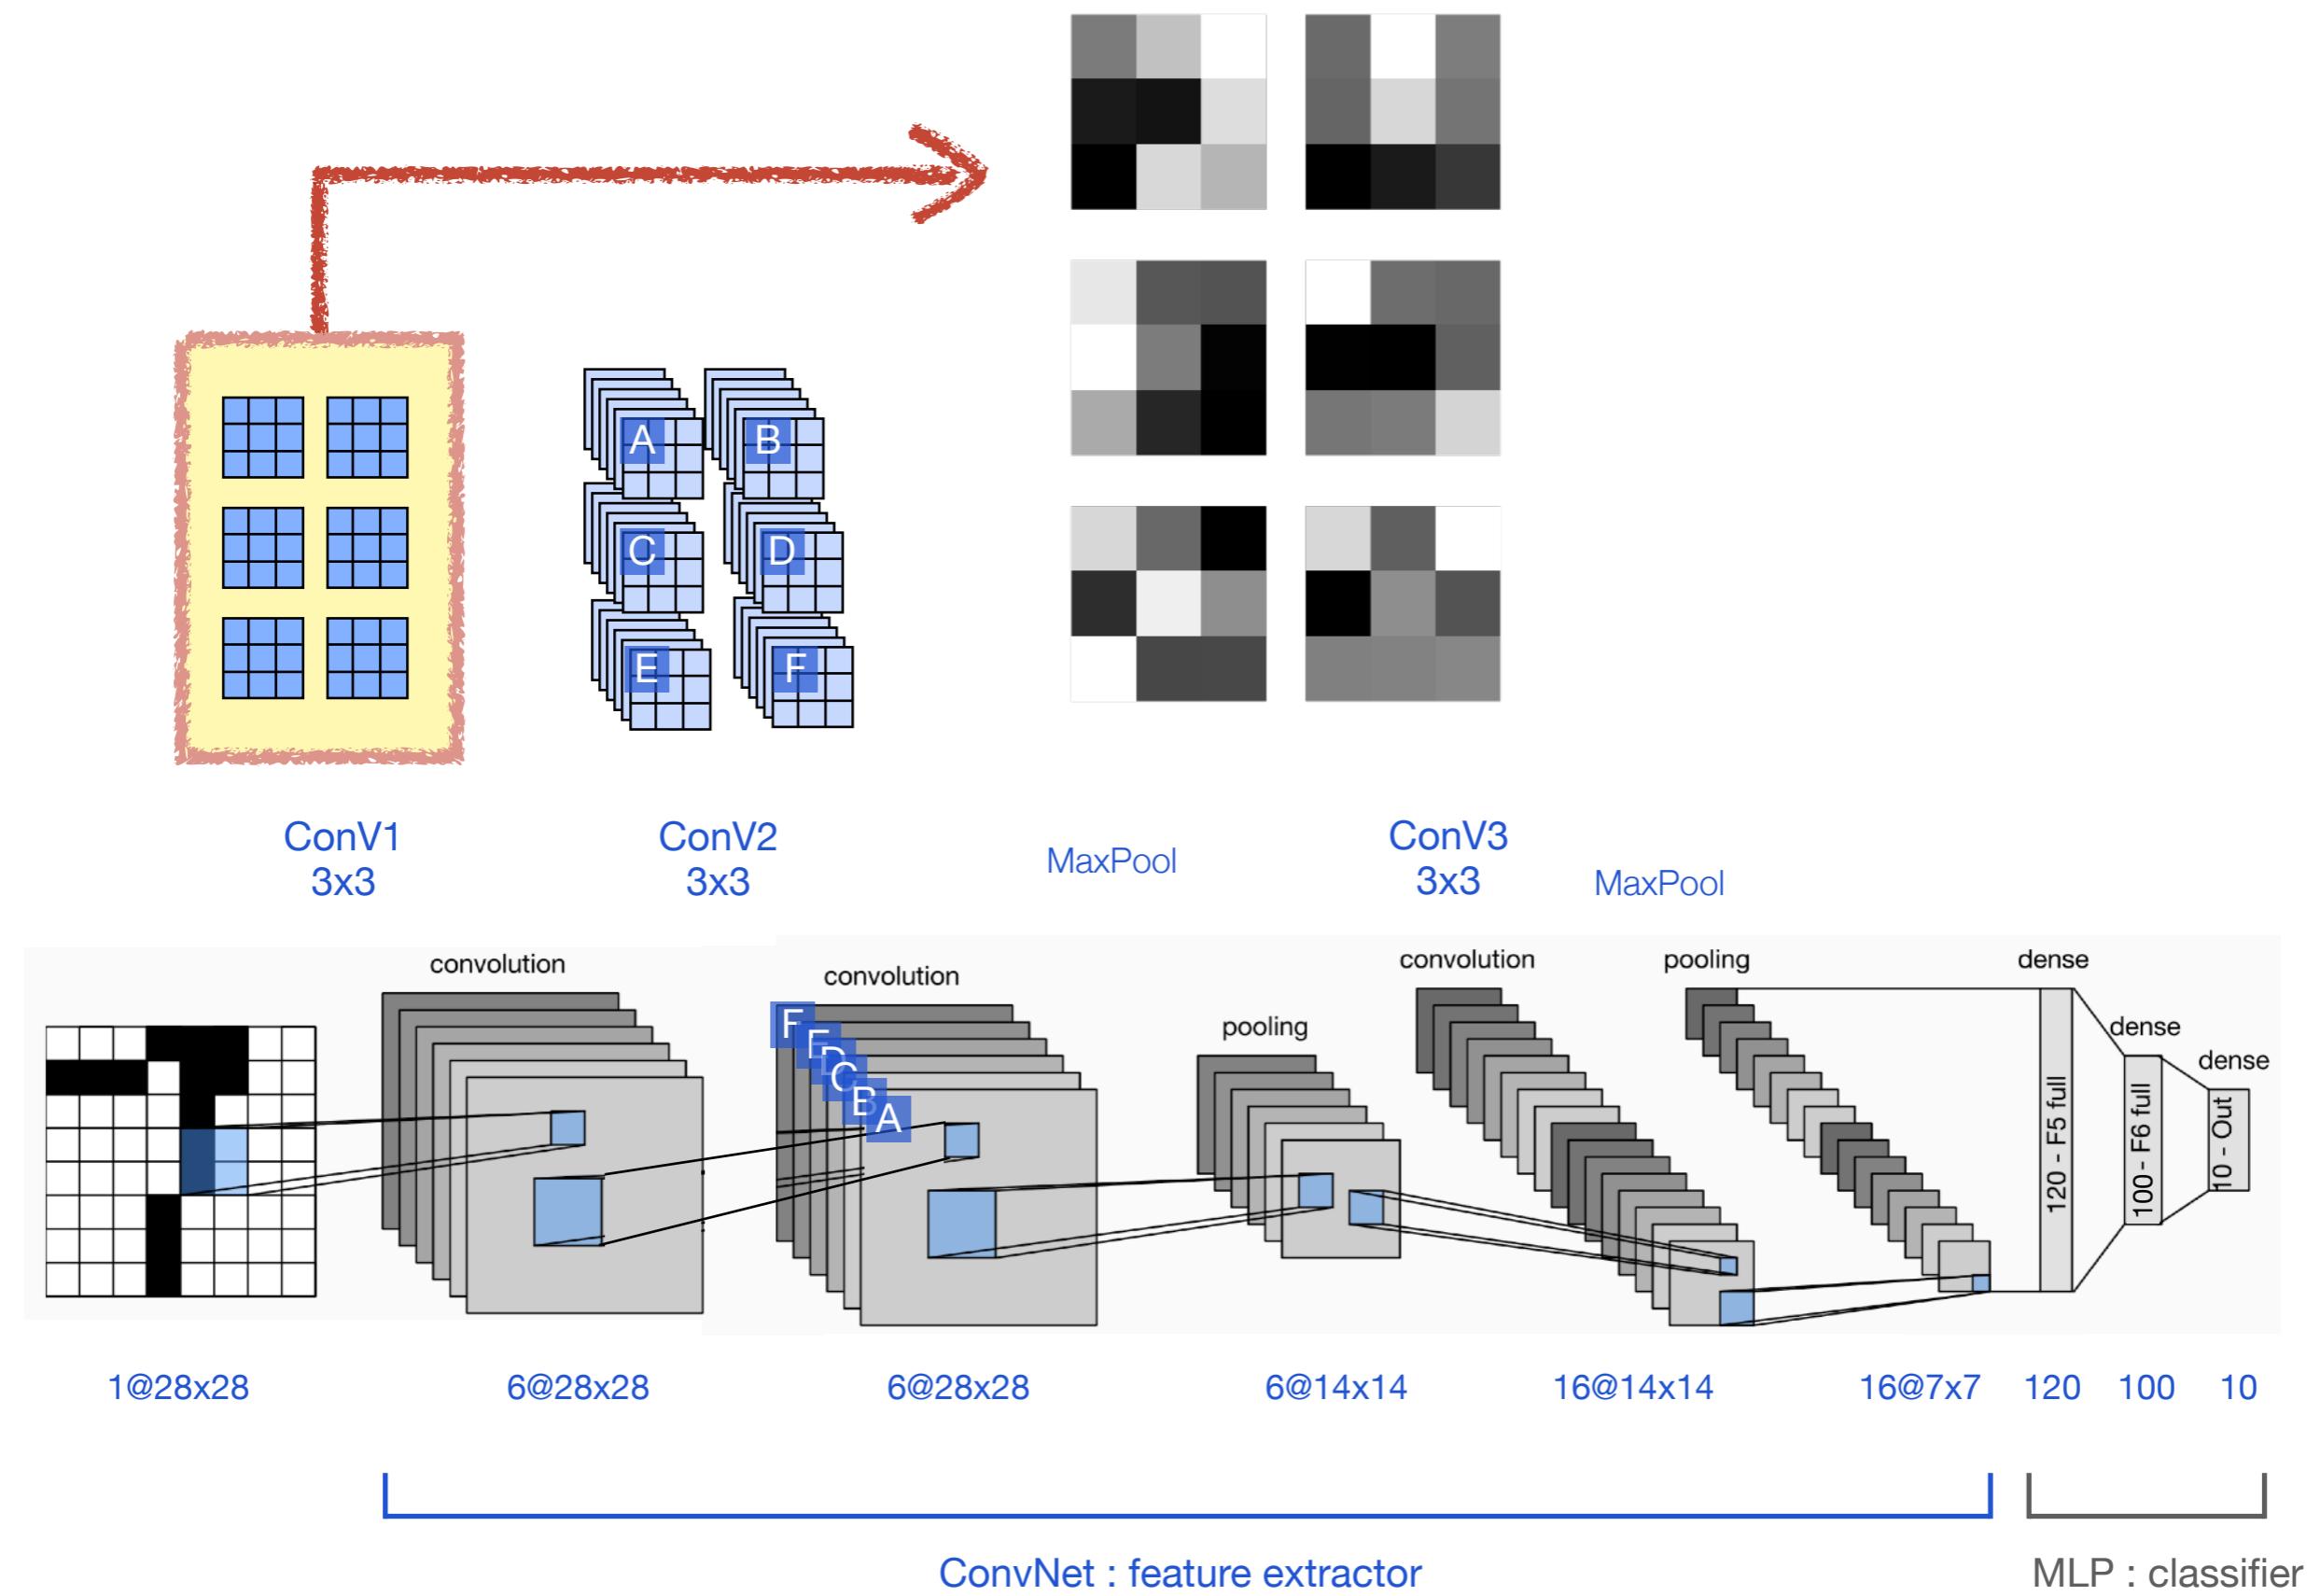
Training

```
# Optimizer
learning_rate = 0.2
optimizer = torch.optim.SGD(net.parameters(), lr=learning_rate)
scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer, milestones=[2, 5, 8, 11, 14], gamma=0.5)
```



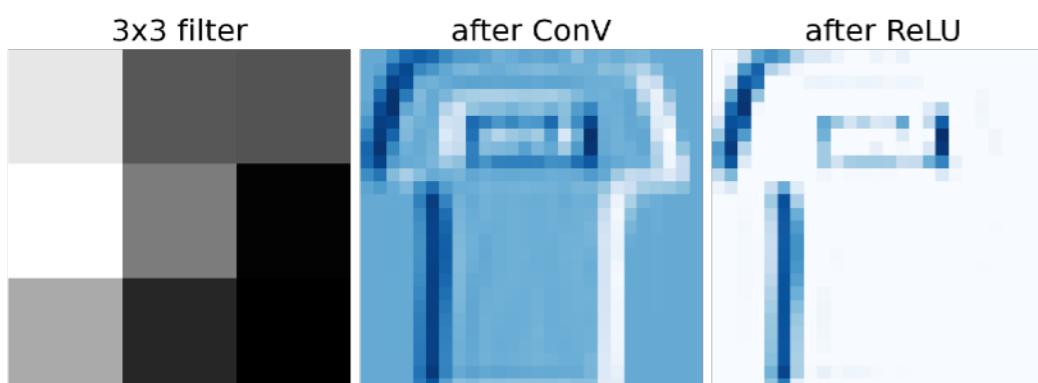
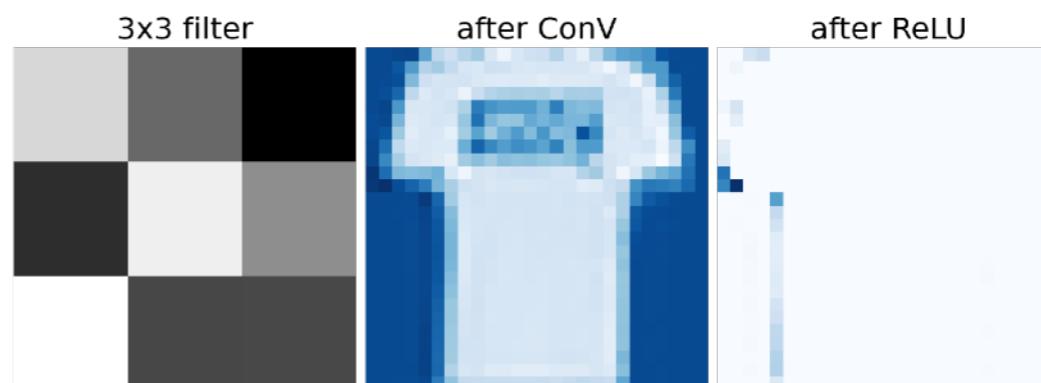
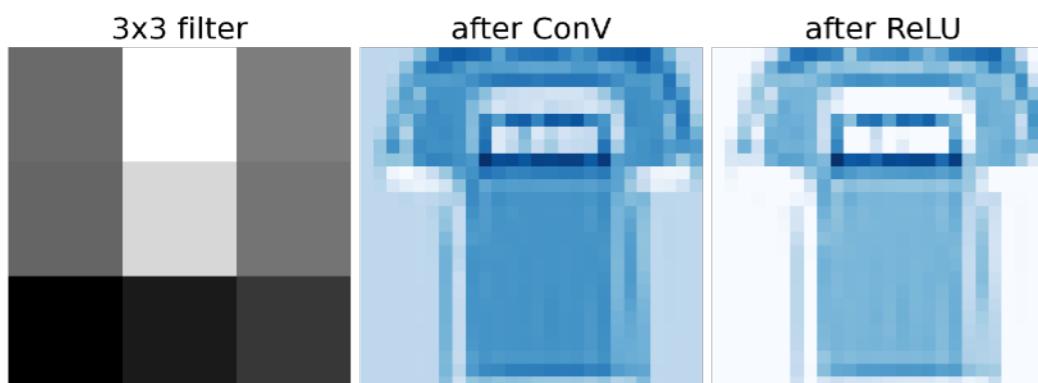
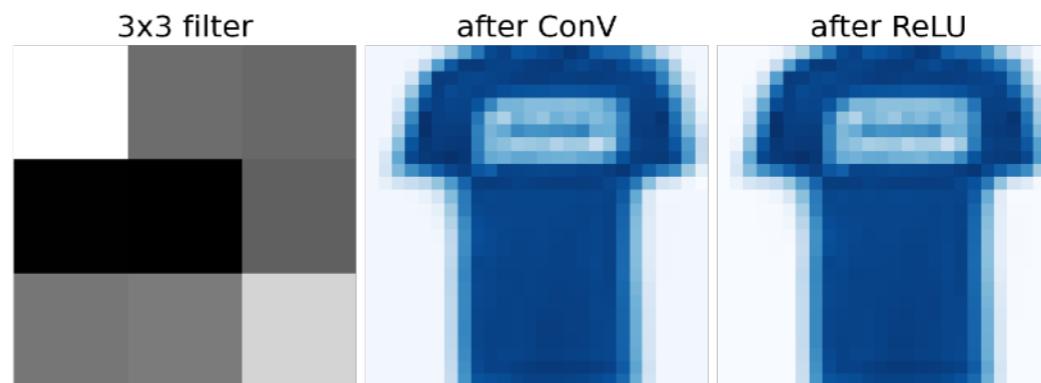
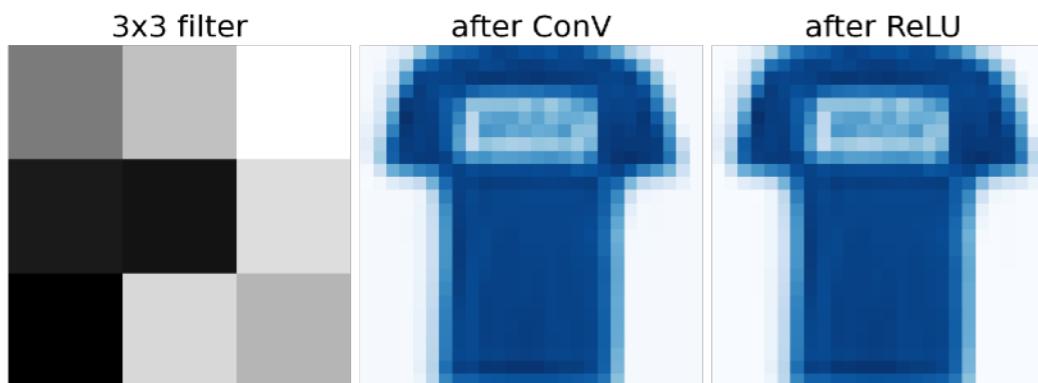
```
>>> End epoch 1, train accuracy = 51.77%, train avgLoss = 1.2684
>>> valid accuracy = 79.38%, valid avgLoss = 0.5493
>>> End epoch 2, train accuracy = 83.00%, train avgLoss = 0.4497
>>> valid accuracy = 83.78%, valid avgLoss = 0.4274
>>> End epoch 3, train accuracy = 87.29%, train avgLoss = 0.3438
>>> valid accuracy = 86.31%, valid avgLoss = 0.3634
>>> End epoch 4, train accuracy = 88.06%, train avgLoss = 0.3190
>>> valid accuracy = 87.02%, valid avgLoss = 0.3569
>>> End epoch 5, train accuracy = 88.74%, train avgLoss = 0.3009
>>> valid accuracy = 87.21%, valid avgLoss = 0.3424
>>> End epoch 6, train accuracy = 89.88%, train avgLoss = 0.2700
>>> valid accuracy = 88.87%, valid avgLoss = 0.3050
>>> End epoch 7, train accuracy = 90.16%, train avgLoss = 0.2609
>>> valid accuracy = 88.85%, valid avgLoss = 0.3059
>>> End epoch 8, train accuracy = 90.49%, train avgLoss = 0.2507
>>> valid accuracy = 89.60%, valid avgLoss = 0.2830
>>> End epoch 9, train accuracy = 91.26%, train avgLoss = 0.2361
>>> valid accuracy = 89.53%, valid avgLoss = 0.2827
>>> End epoch 10, train accuracy = 91.39%, train avgLoss = 0.2317
>>> valid accuracy = 89.70%, valid avgLoss = 0.2818
>>> End epoch 11, train accuracy = 91.48%, train avgLoss = 0.2283
>>> valid accuracy = 89.59%, valid avgLoss = 0.2839
>>> End epoch 12, train accuracy = 91.86%, train avgLoss = 0.2188
>>> valid accuracy = 90.00%, valid avgLoss = 0.2739
>>> End epoch 13, train accuracy = 91.97%, train avgLoss = 0.2167
>>> valid accuracy = 90.13%, valid avgLoss = 0.2684
>>> End epoch 14, train accuracy = 92.05%, train avgLoss = 0.2143
>>> valid accuracy = 90.28%, valid avgLoss = 0.2687
>>> End epoch 15, train accuracy = 92.17%, train avgLoss = 0.2093
>>> valid accuracy = 90.39%, valid avgLoss = 0.2662
```

1st layer convolutional filters



original

1st layer convolutional filters



1st layer convolutional filters of AlexNet trained on the ImageNet dataset

- Load pre-trained AlexNet in Pytorch from `torchvision.models`

ImageNet dataset
~ 1.3 million images to train
50,000 images for validation
1000 categories

```
Input shape: torch.Size([1, 3, 224, 224])
Conv2d output shape: torch.Size([1, 64, 55, 55])
ReLU output shape: torch.Size([1, 64, 55, 55])
MaxPool2d output shape: torch.Size([1, 64, 27, 27])
Conv2d output shape: torch.Size([1, 192, 27, 27])
ReLU output shape: torch.Size([1, 192, 27, 27])
MaxPool2d output shape: torch.Size([1, 192, 13, 13])
Conv2d output shape: torch.Size([1, 384, 13, 13])
ReLU output shape: torch.Size([1, 384, 13, 13])
Conv2d output shape: torch.Size([1, 256, 13, 13])
ReLU output shape: torch.Size([1, 256, 13, 13])
Conv2d output shape: torch.Size([1, 256, 13, 13])
ReLU output shape: torch.Size([1, 256, 13, 13])
MaxPool2d output shape: torch.Size([1, 256, 6, 6])
```

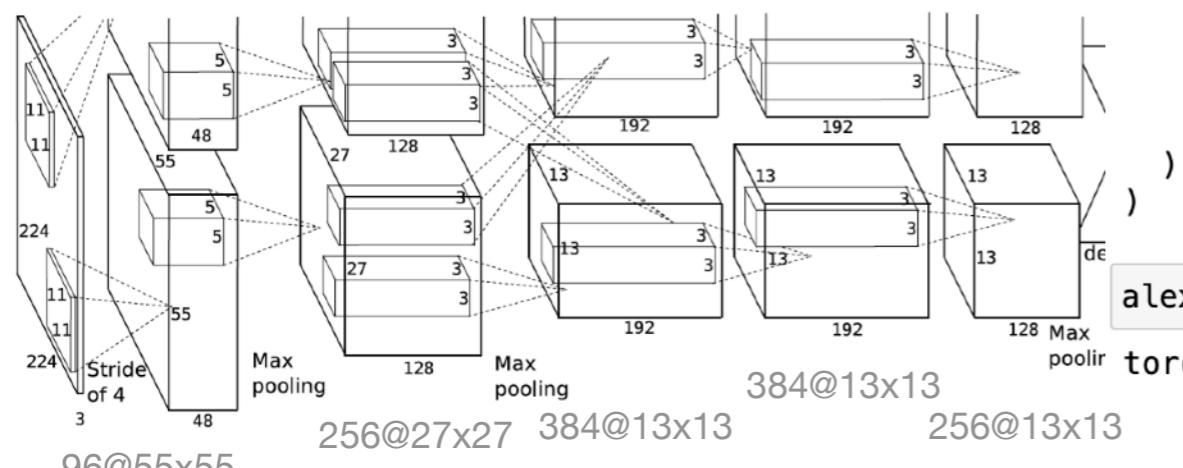
```
import torchvision.models as models
alexnet = models.alexnet(pretrained=True)
```

```
alexnet
```

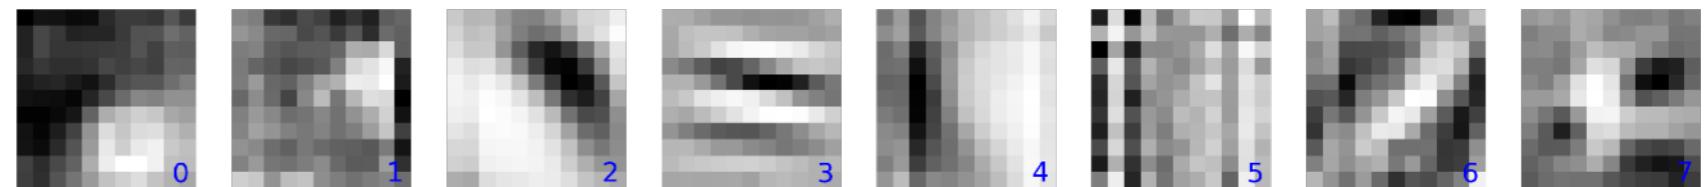
```
AlexNet(
    (features): Sequential(
        (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
        (1): ReLU(inplace=True)
        (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
        (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
        (4): ReLU(inplace=True)
        (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
        (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (7): ReLU(inplace=True)
        (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (9): ReLU(inplace=True)
        (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (11): ReLU(inplace=True)
        (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
    (classifier): Sequential(
        (0): Dropout(p=0.5, inplace=False)
        (1): Linear(in_features=9216, out_features=4096, bias=True)
        (2): ReLU(inplace=True)
        (3): Dropout(p=0.5, inplace=False)
        (4): Linear(in_features=4096, out_features=4096, bias=True)
        (5): ReLU(inplace=True)
        (6): Linear(in_features=4096, out_features=1000, bias=True)
    )
)
```

```
alexnet.features[0].weight.data.shape
```

```
torch.Size([64, 3, 11, 11])
```



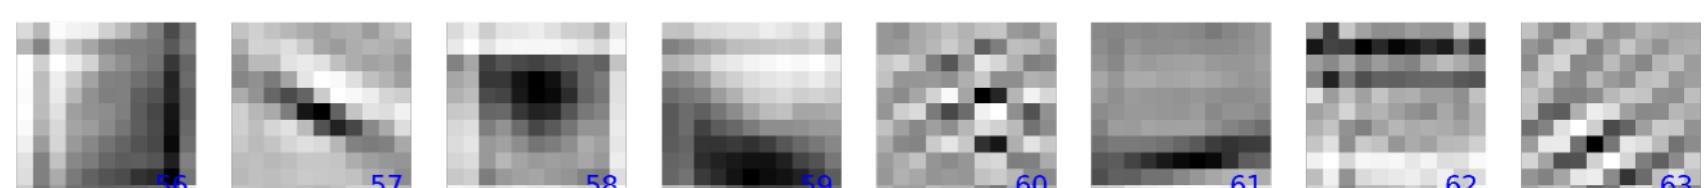
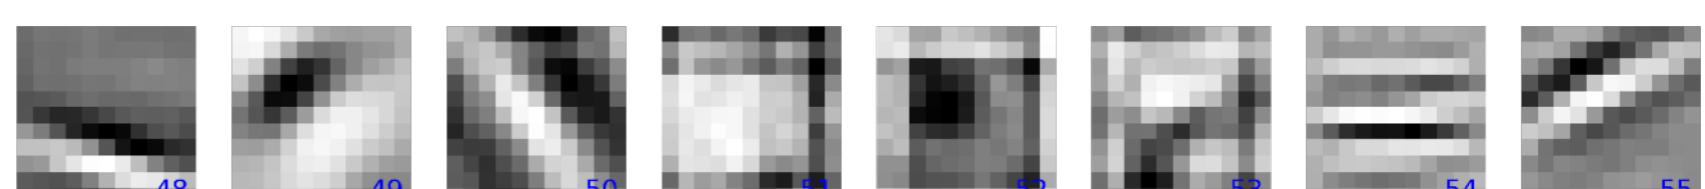
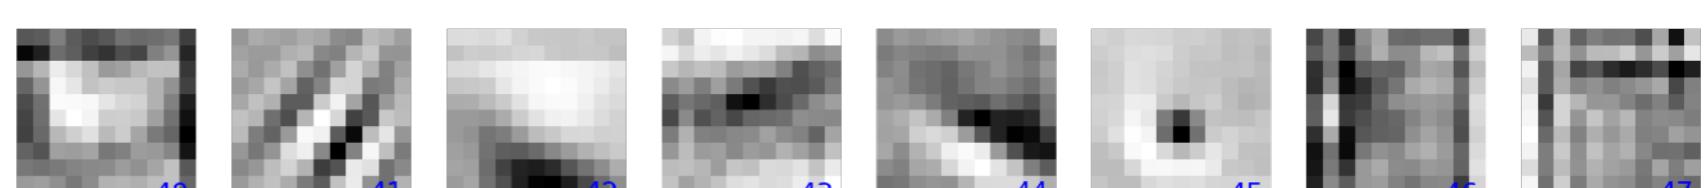
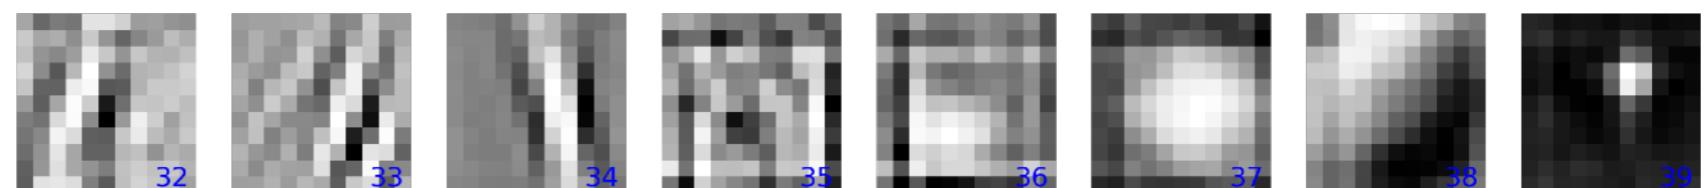
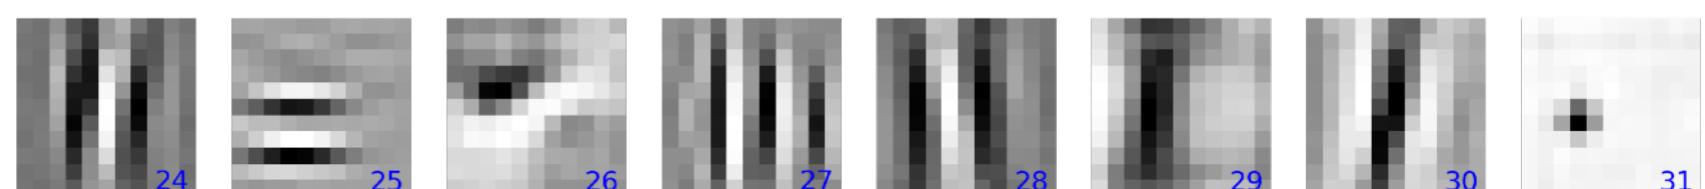
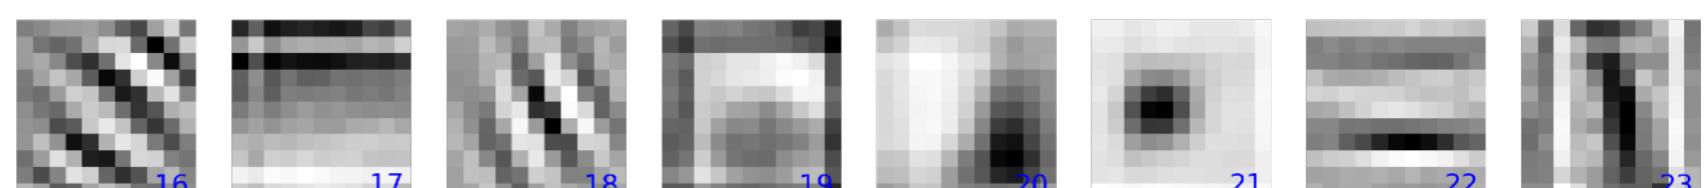
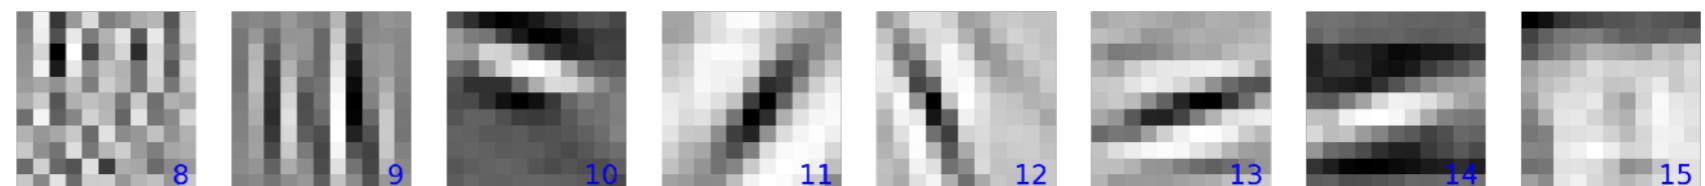
1st layer convolutional filters of AlexNet trained on ImageNet dataset



1st layer ConV weight matrix :

64 x 3 x 11 x 11

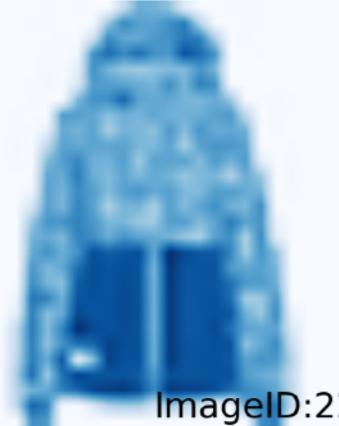
kernel size



```
Input shape:      torch.Size([1, 3, 224, 224])
Conv2d output shape:  torch.Size([1, 64, 55, 55])
ReLU output shape:   torch.Size([1, 64, 55, 55])
MaxPool2d output shape: torch.Size([1, 64, 27, 27])
Conv2d output shape:  torch.Size([1, 192, 27, 27])
ReLU output shape:   torch.Size([1, 192, 27, 27])
MaxPool2d output shape: torch.Size([1, 192, 13, 13])
Conv2d output shape:  torch.Size([1, 384, 13, 13])
ReLU output shape:   torch.Size([1, 384, 13, 13])
Conv2d output shape:  torch.Size([1, 256, 13, 13])
ReLU output shape:   torch.Size([1, 256, 13, 13])
Conv2d output shape:  torch.Size([1, 256, 13, 13])
ReLU output shape:   torch.Size([1, 256, 13, 13])
MaxPool2d output shape: torch.Size([1, 256, 6, 6])
```

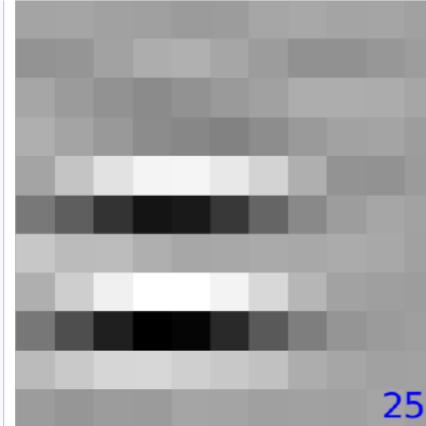
1st layer convolutional filters of AlexNet trained on ImageNet dataset and applied on Fashion-MNIST

original



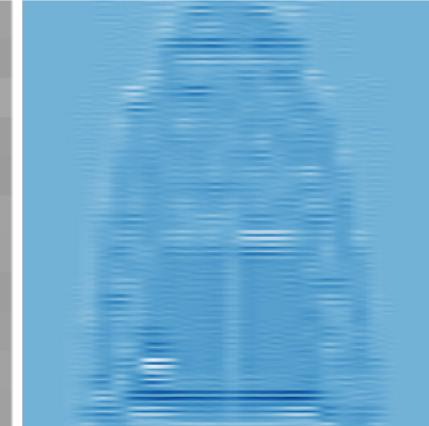
ImageID:22

11x11 ConV



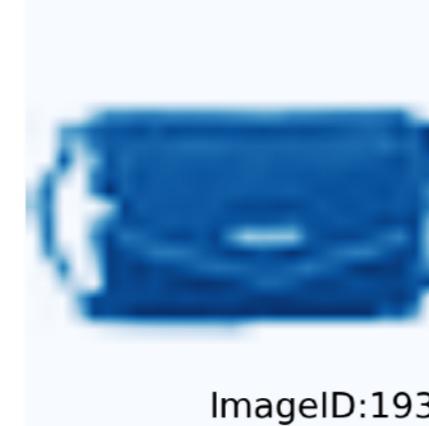
25

feature map



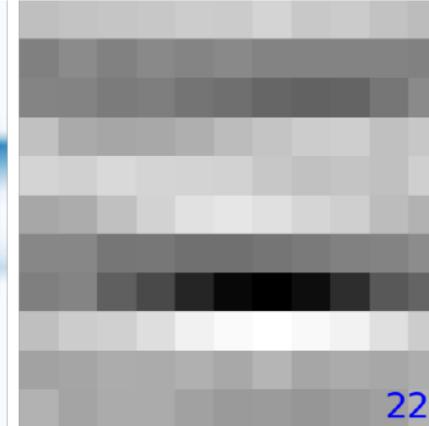
25

original



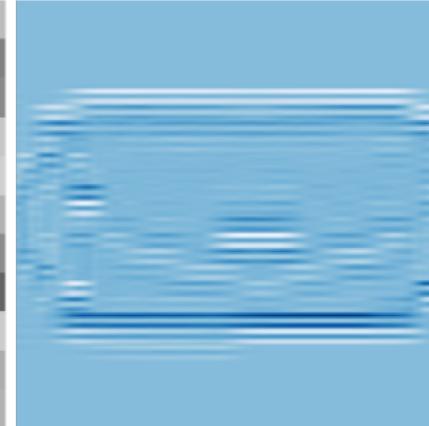
ImageID:193

11x11 ConV



22

feature map

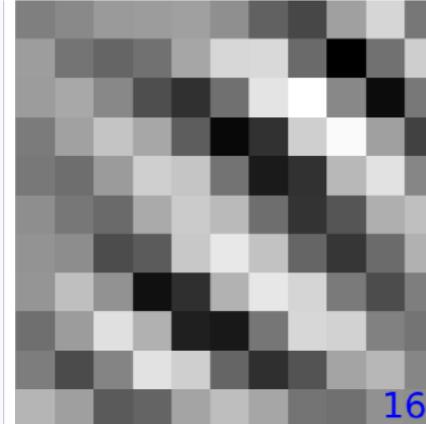


original



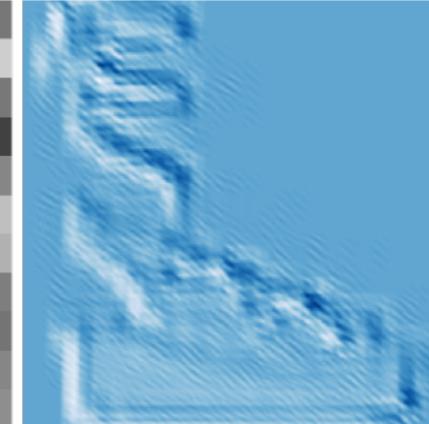
ImageID:36

11x11 ConV



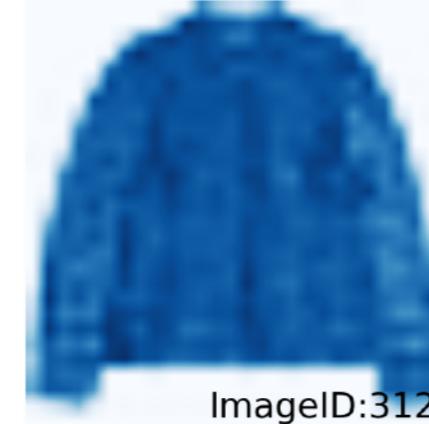
16

feature map



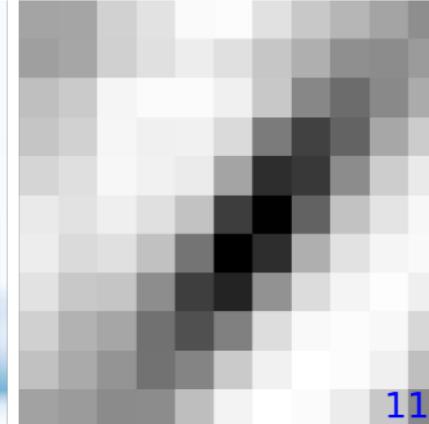
16

original



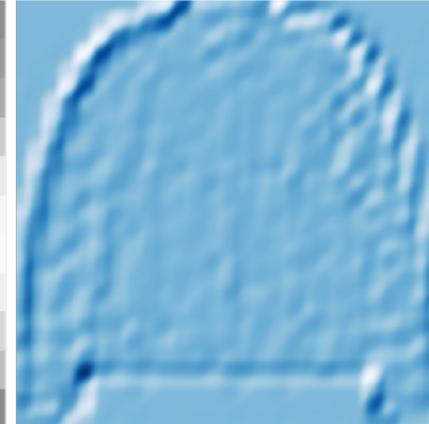
ImageID:312

11x11 ConV



11

feature map

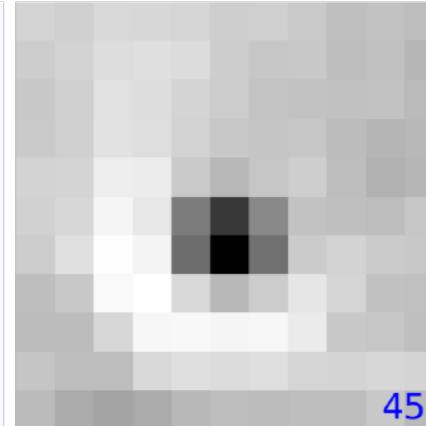


original



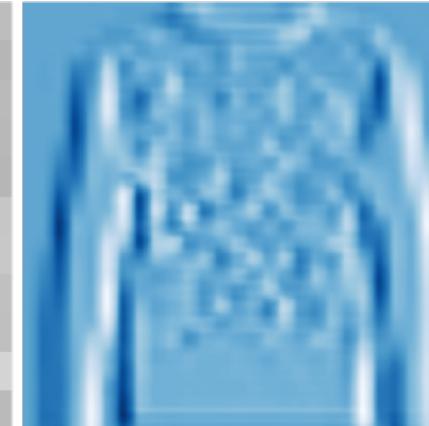
ImageID:1087

11x11 ConV



45

feature map



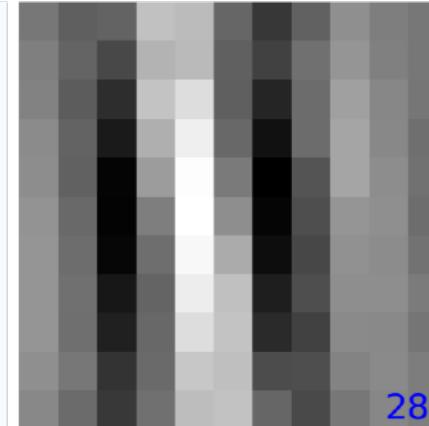
45

original



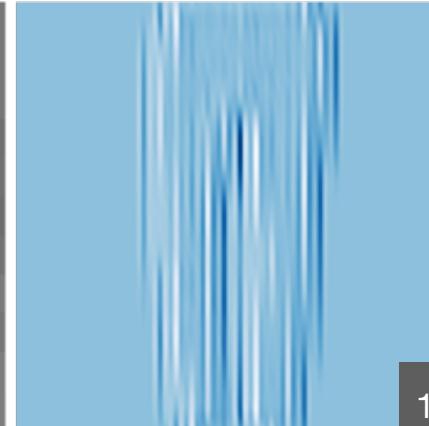
ImageID:74

11x11 ConV



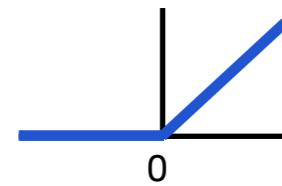
28

feature map



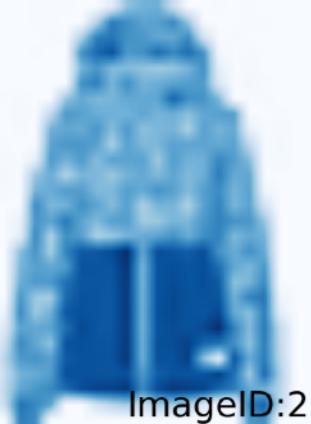
10

1st layer convolutional filters of AlexNet trained on ImageNet dataset and applied on Fashion-MNIST



after applying **ReLU** activation ...

original

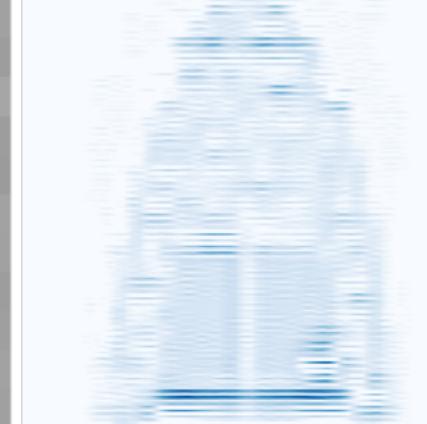


ImageID:22

11x11 ConV



feature map



original

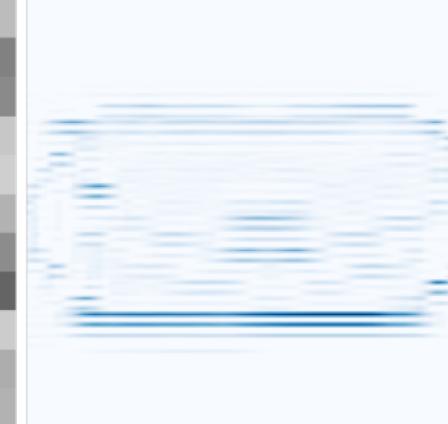


ImageID:193

11x11 ConV



feature map

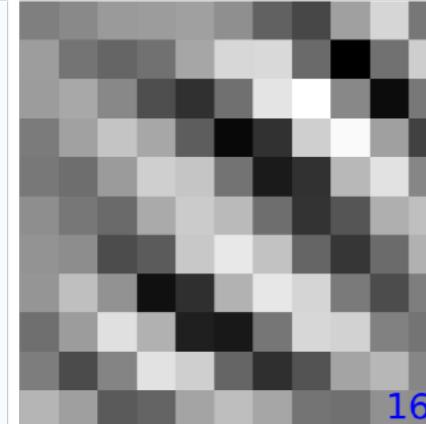


original



ImageID:36

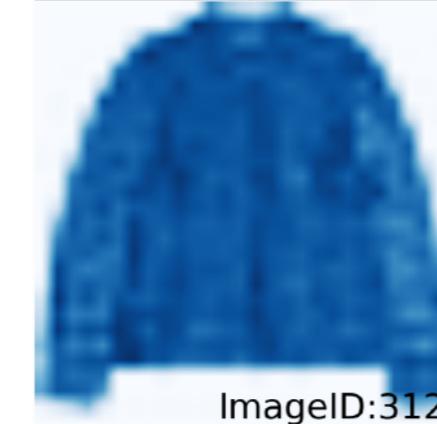
11x11 ConV



feature map

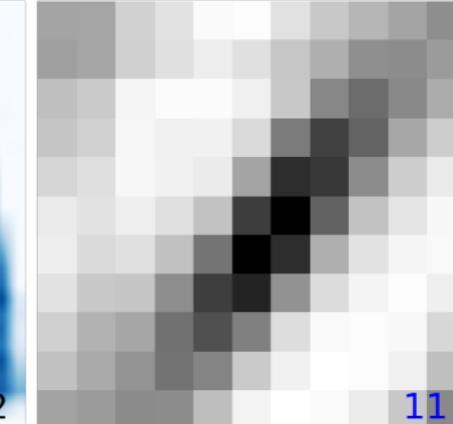


original

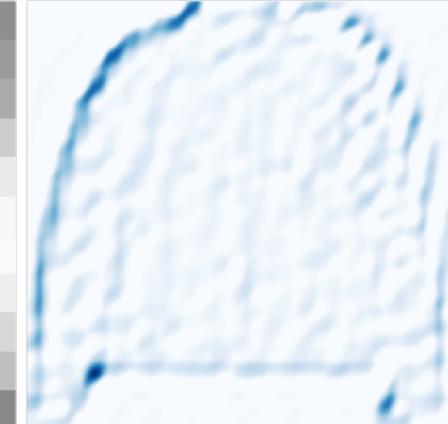


ImageID:312

11x11 ConV



feature map

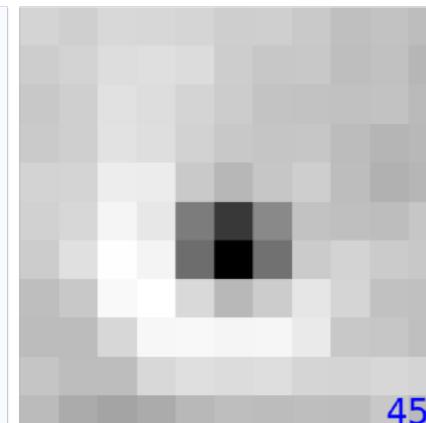


original

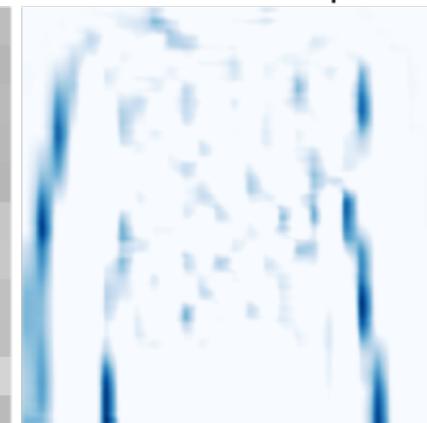


ImageID:1087

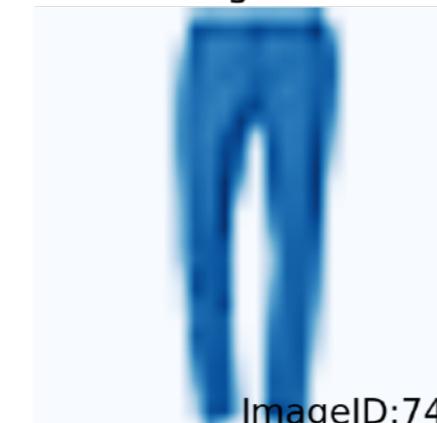
11x11 ConV



feature map

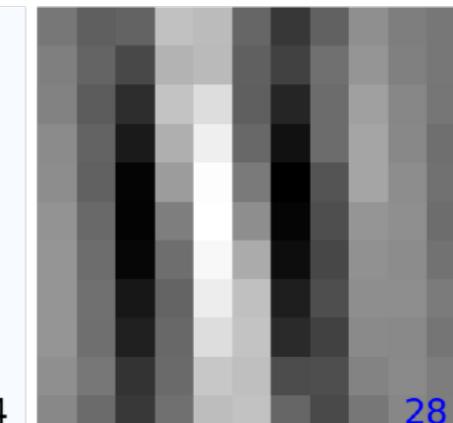


original

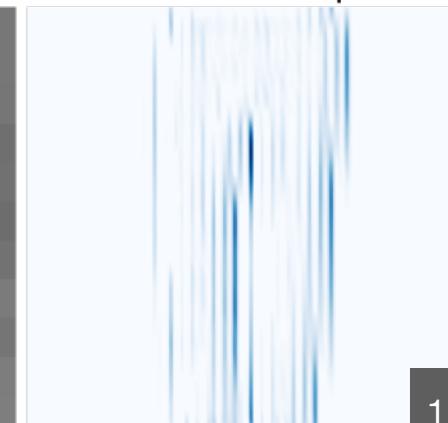


ImageID:74

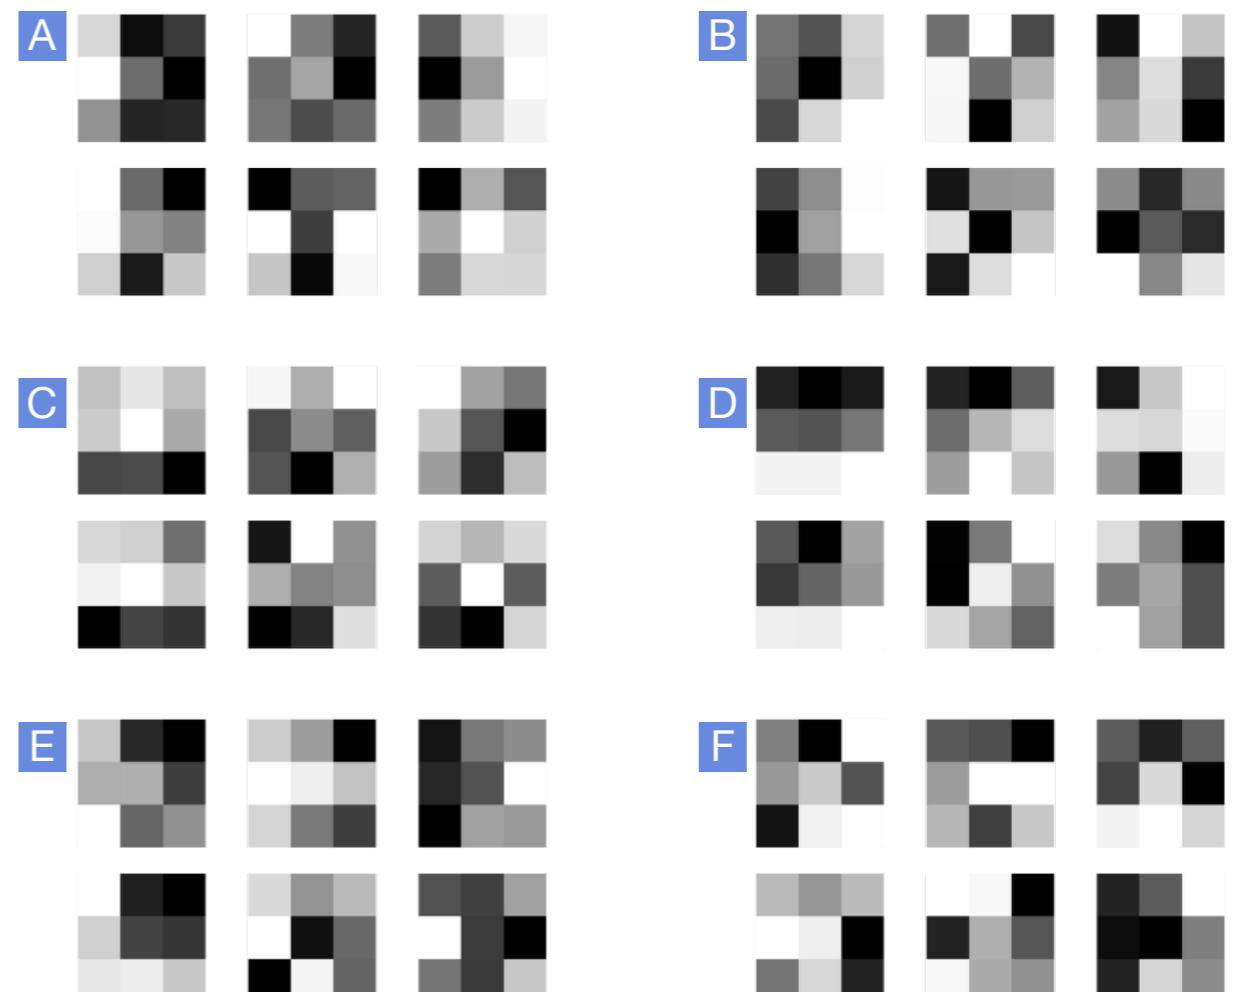
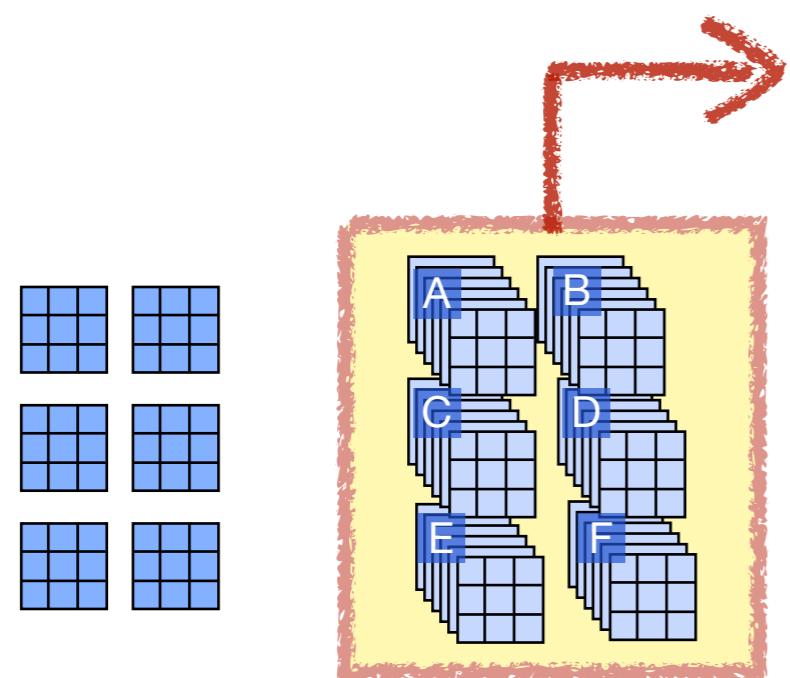
11x11 ConV



feature map



2nd layer convolutional filters



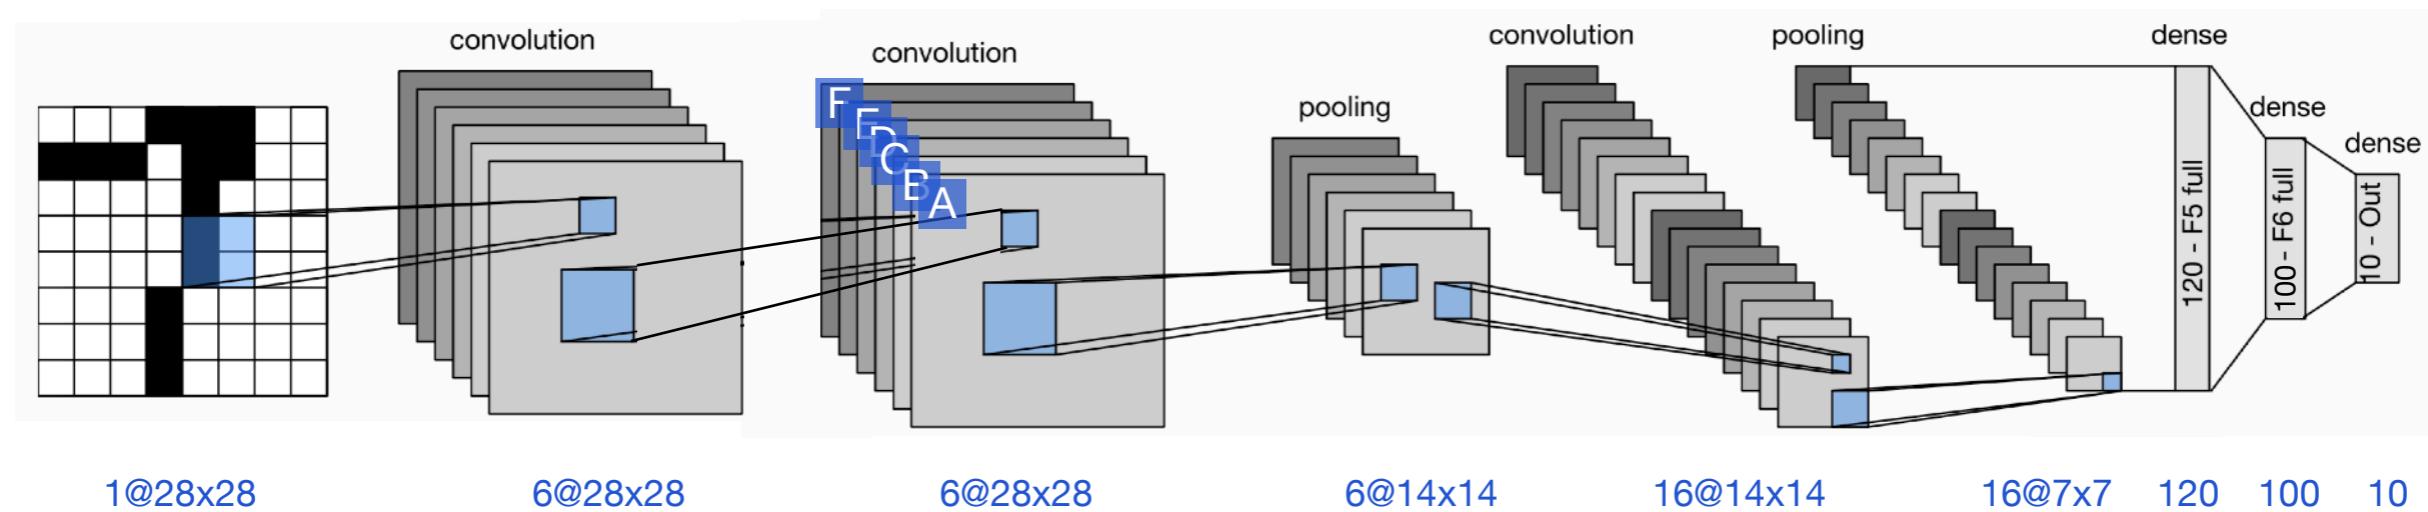
ConV1
3x3

ConV2
3x3

MaxPool

ConV3
3x3

MaxPool



1@28x28

6@28x28

6@28x28

16@14x14

16@7x7

120

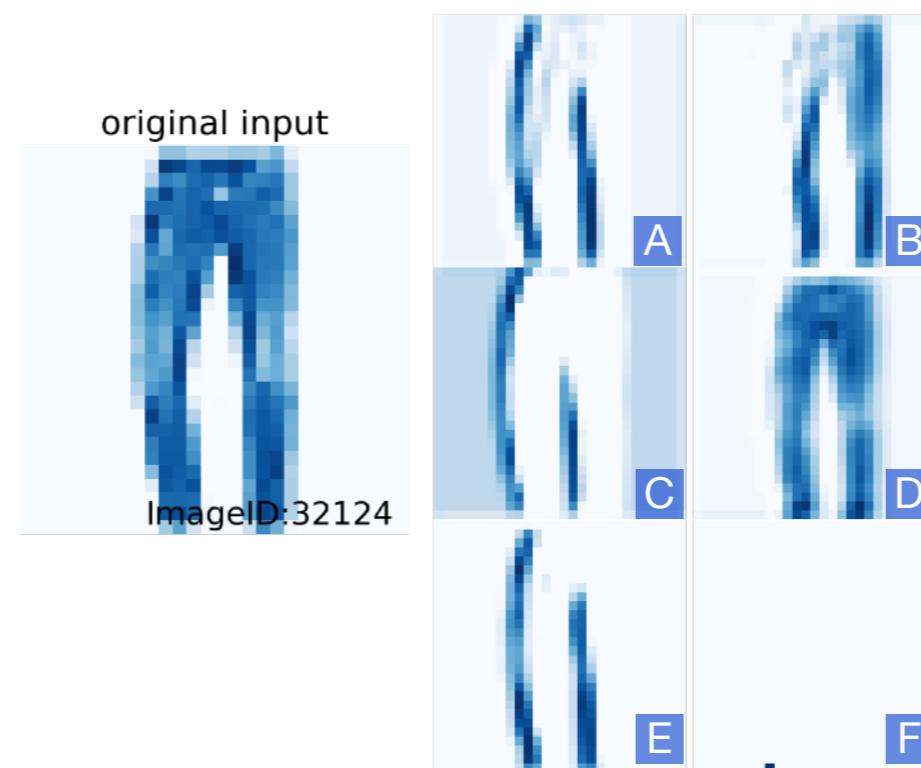
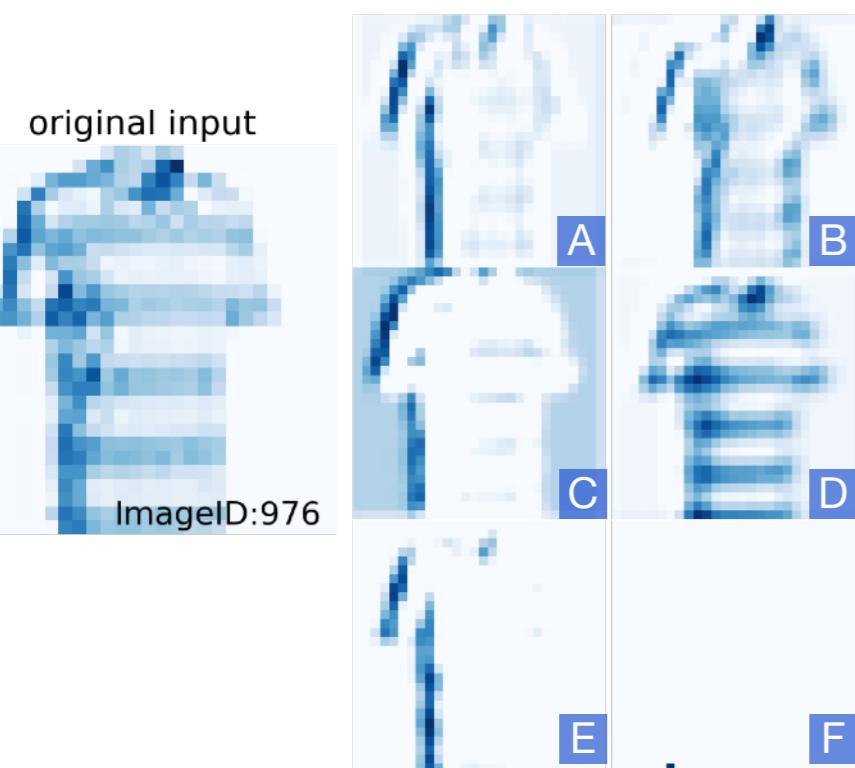
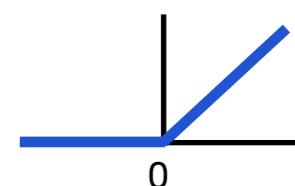
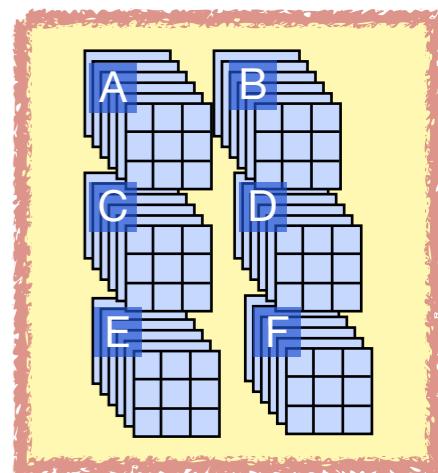
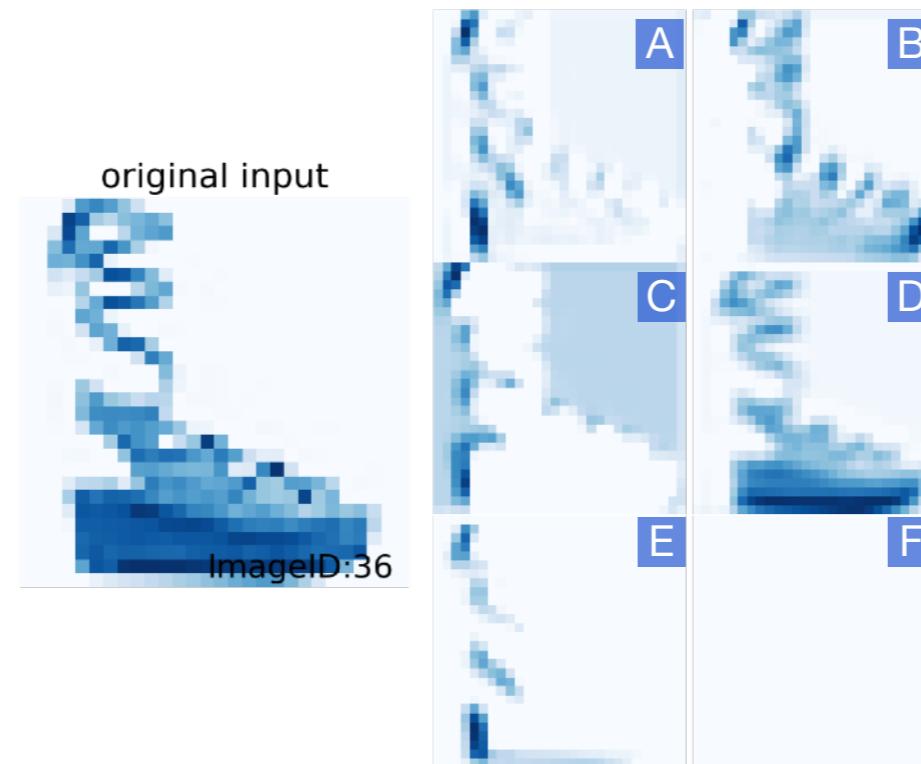
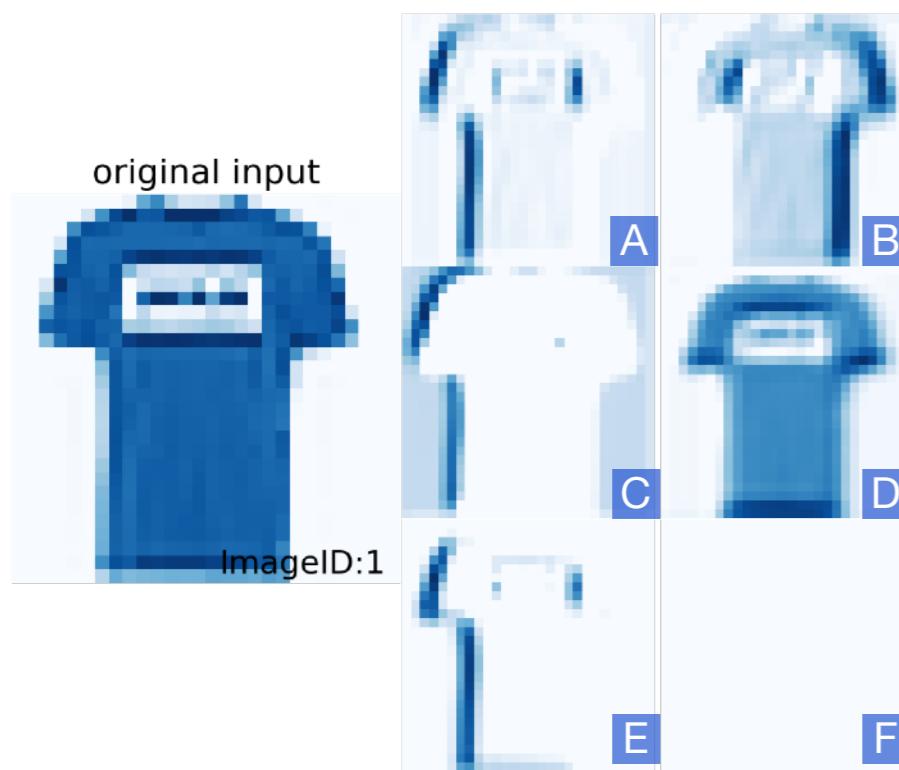
100

10

ConvNet : feature extractor

MLP : classifier

Feature Maps out from ConV2 filters + ReLU



Dying ReLU Problem

When ReLU neurons become inactive and only output 0 for any input.

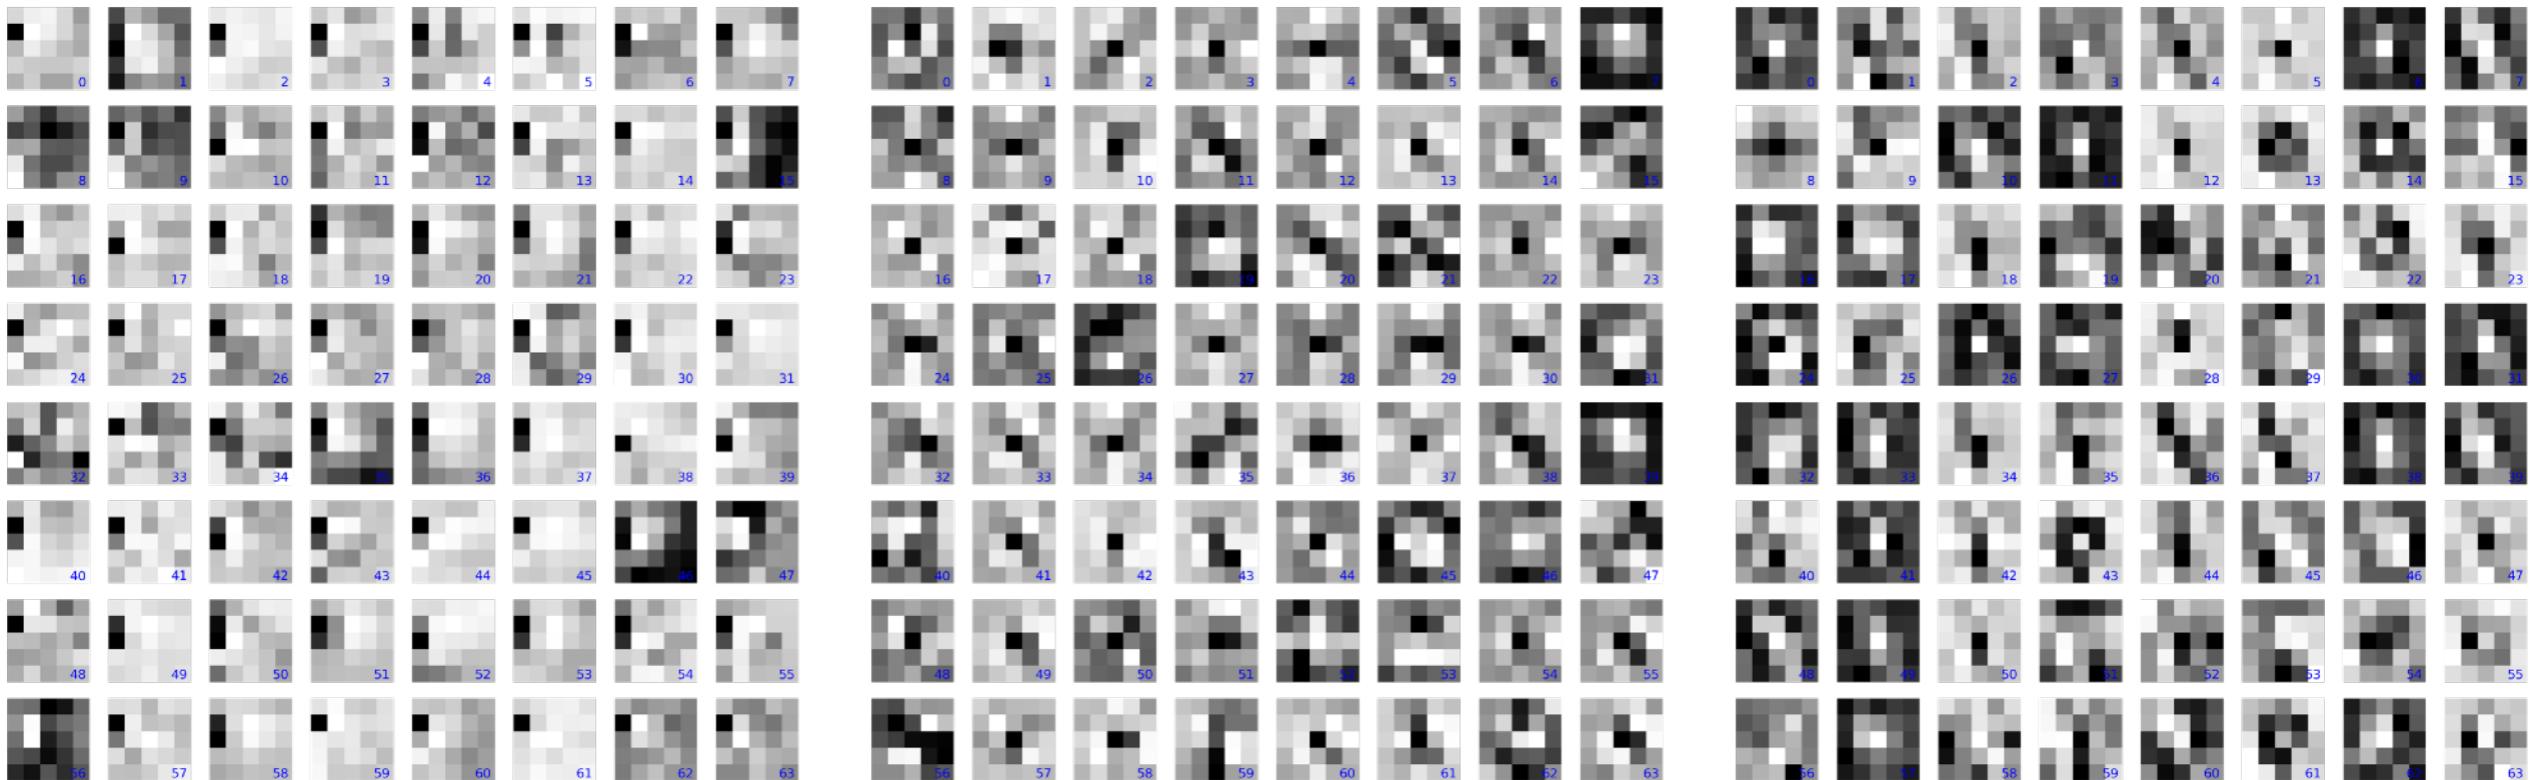
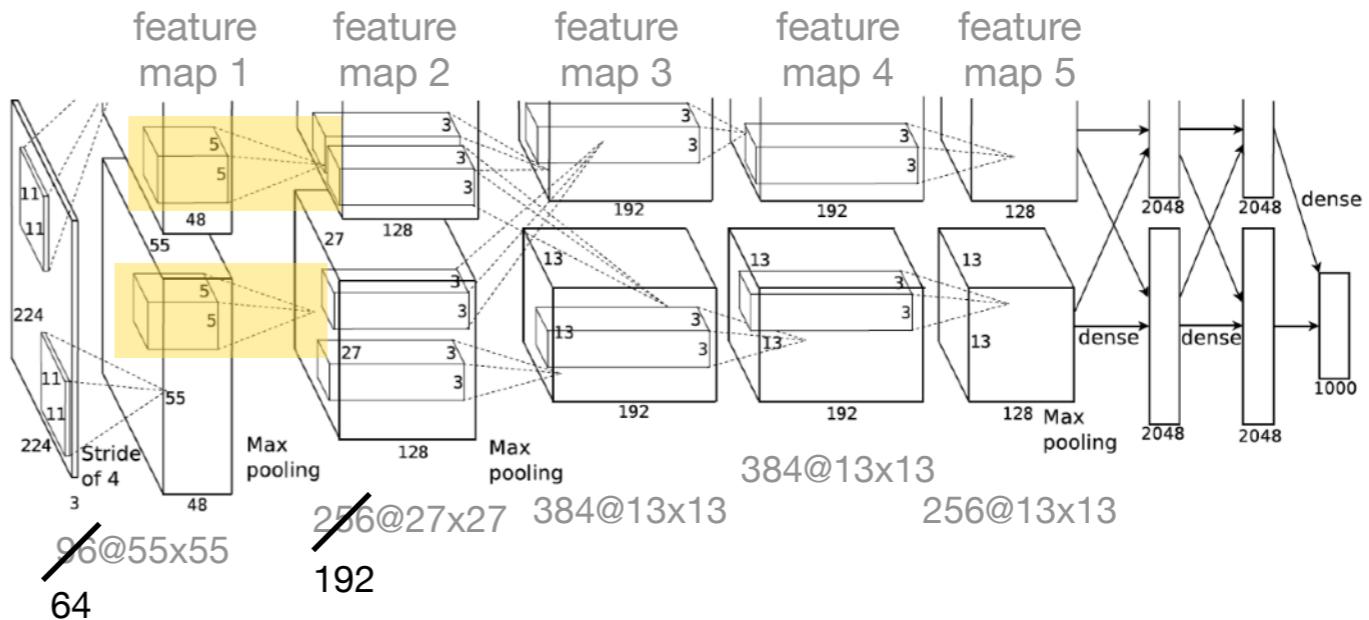
Happens when

1. learning rate is too high.
2. large negative bias

$$\mathbf{w} := \mathbf{w} - \eta \nabla_{\mathbf{w}} L_{\text{mini}}(\mathbf{w})$$

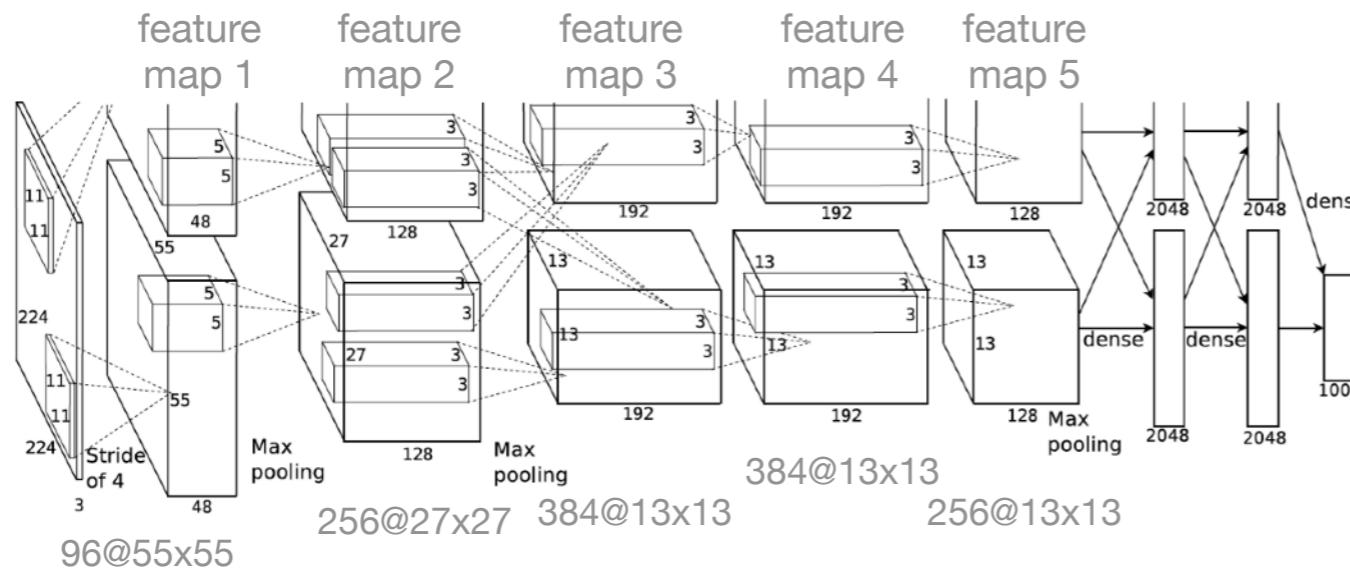
2nd layer convolutional filters of AlexNet trained on ImageNet dataset

2nd layer ConV weight matrix :
192 x 64 x 5 x 5
kernel size

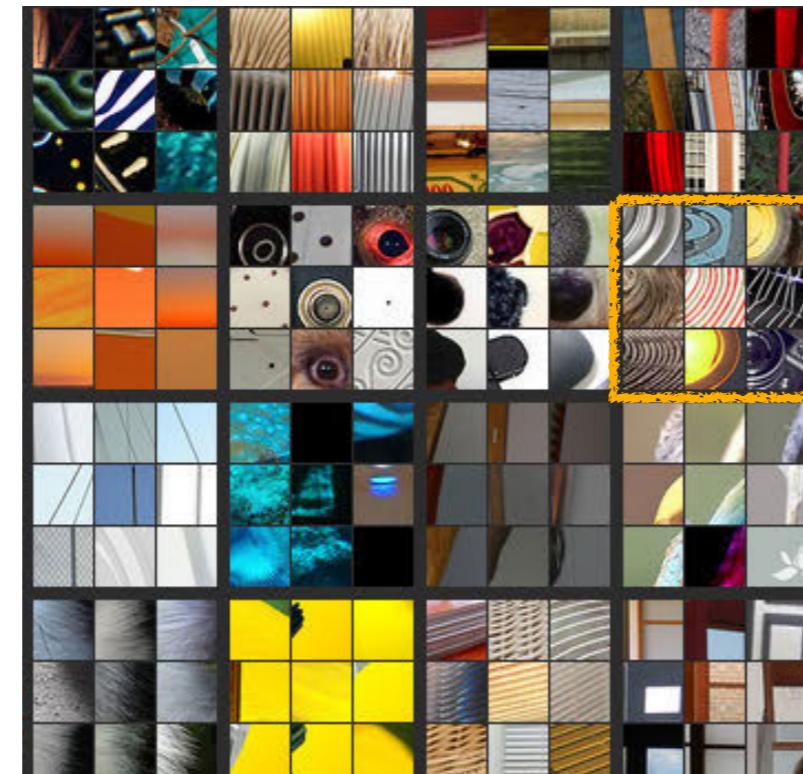


... and another 189 of the 64x5x5 filters

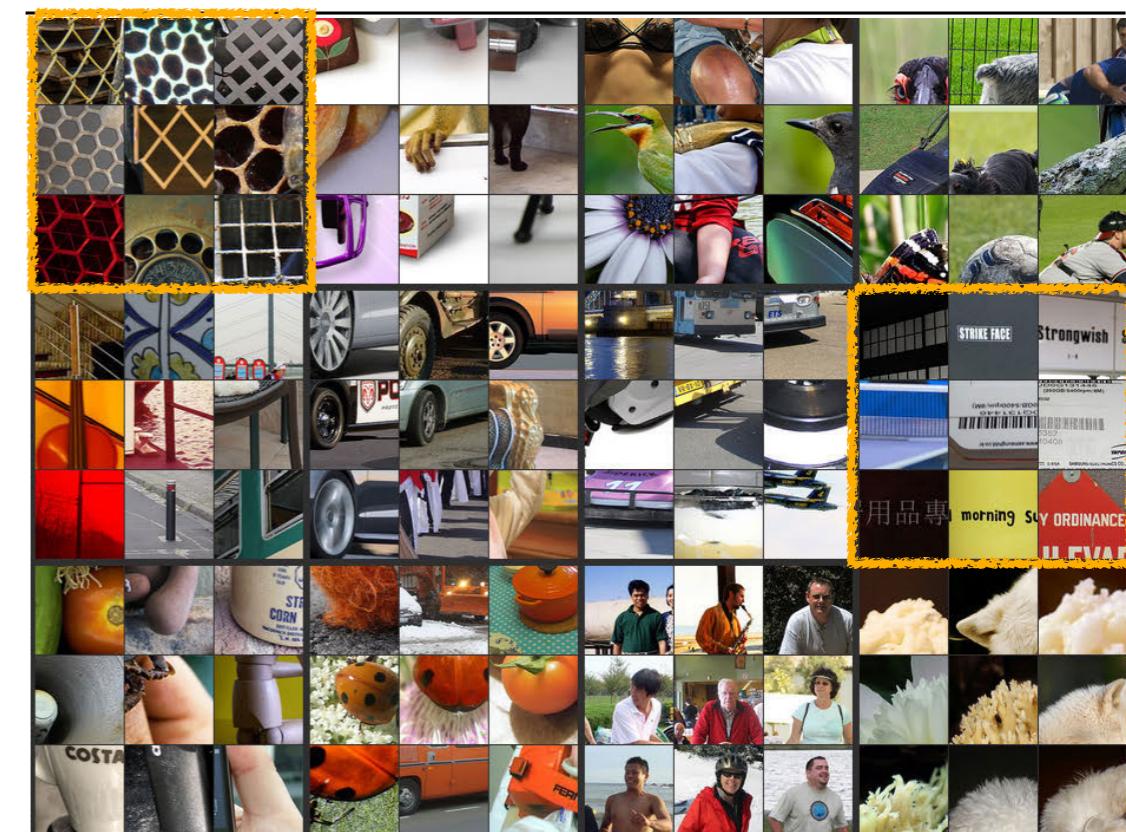
Maximally activated image patches (for AlexNet trained on ImageNet)



layer 1



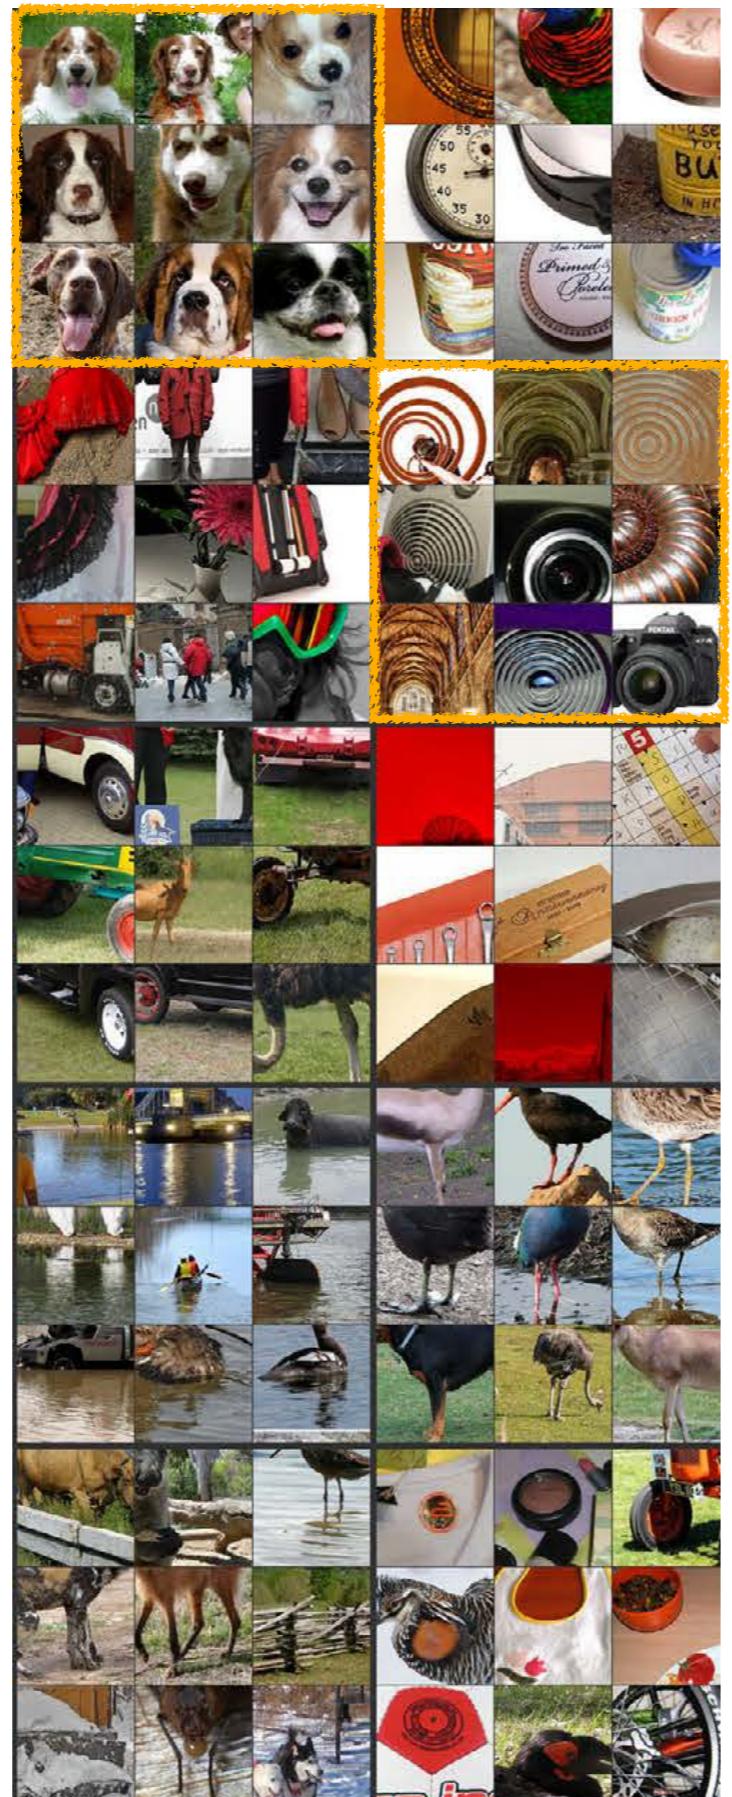
layer 2



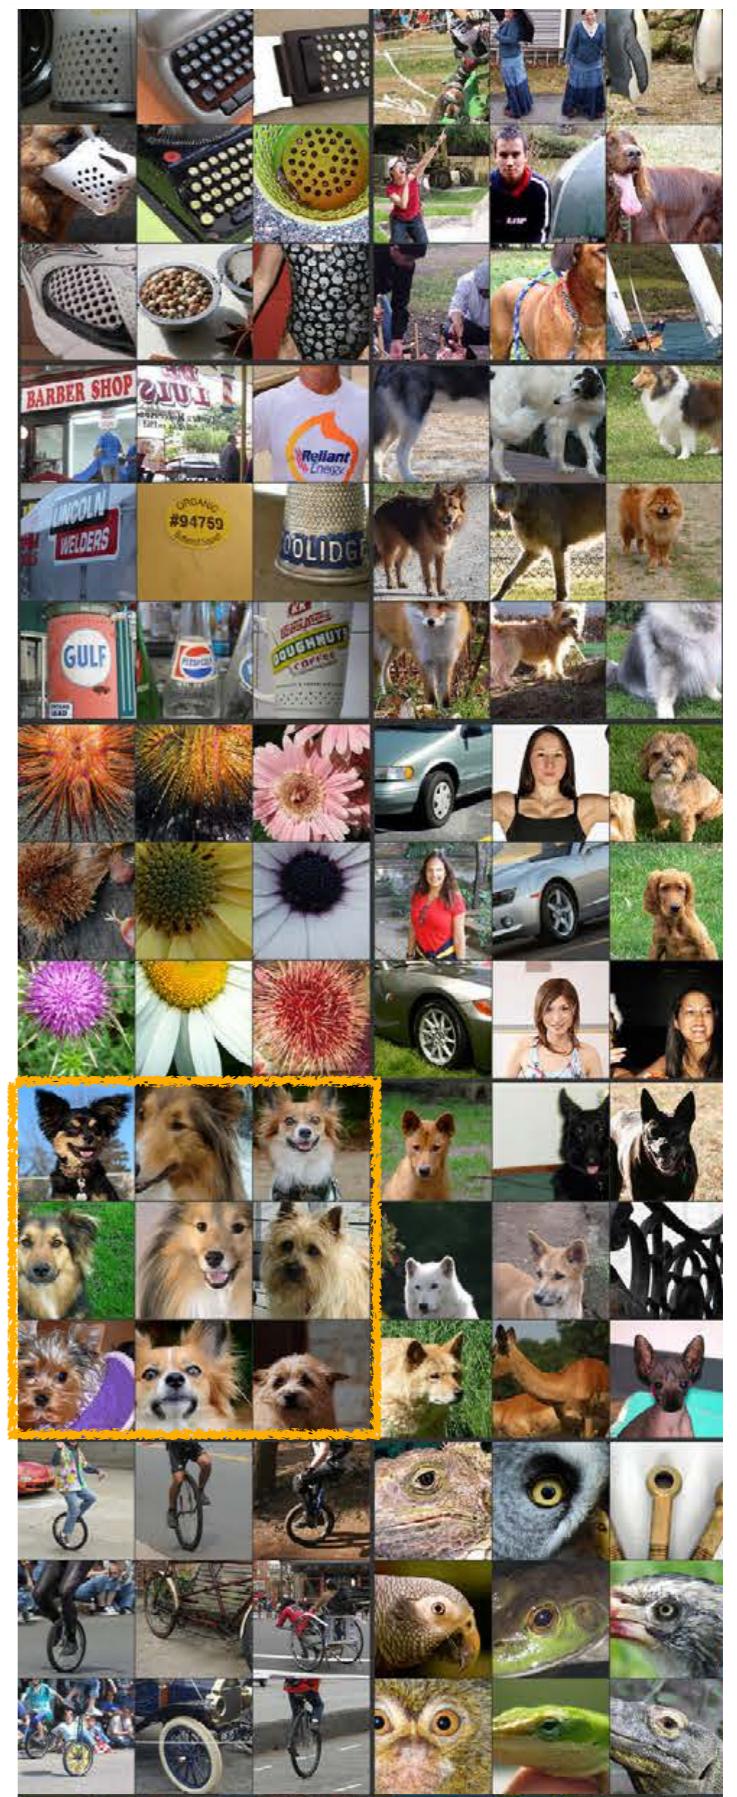
layer 3

Maximally activated image patches (for AlexNet trained on ImageNet)

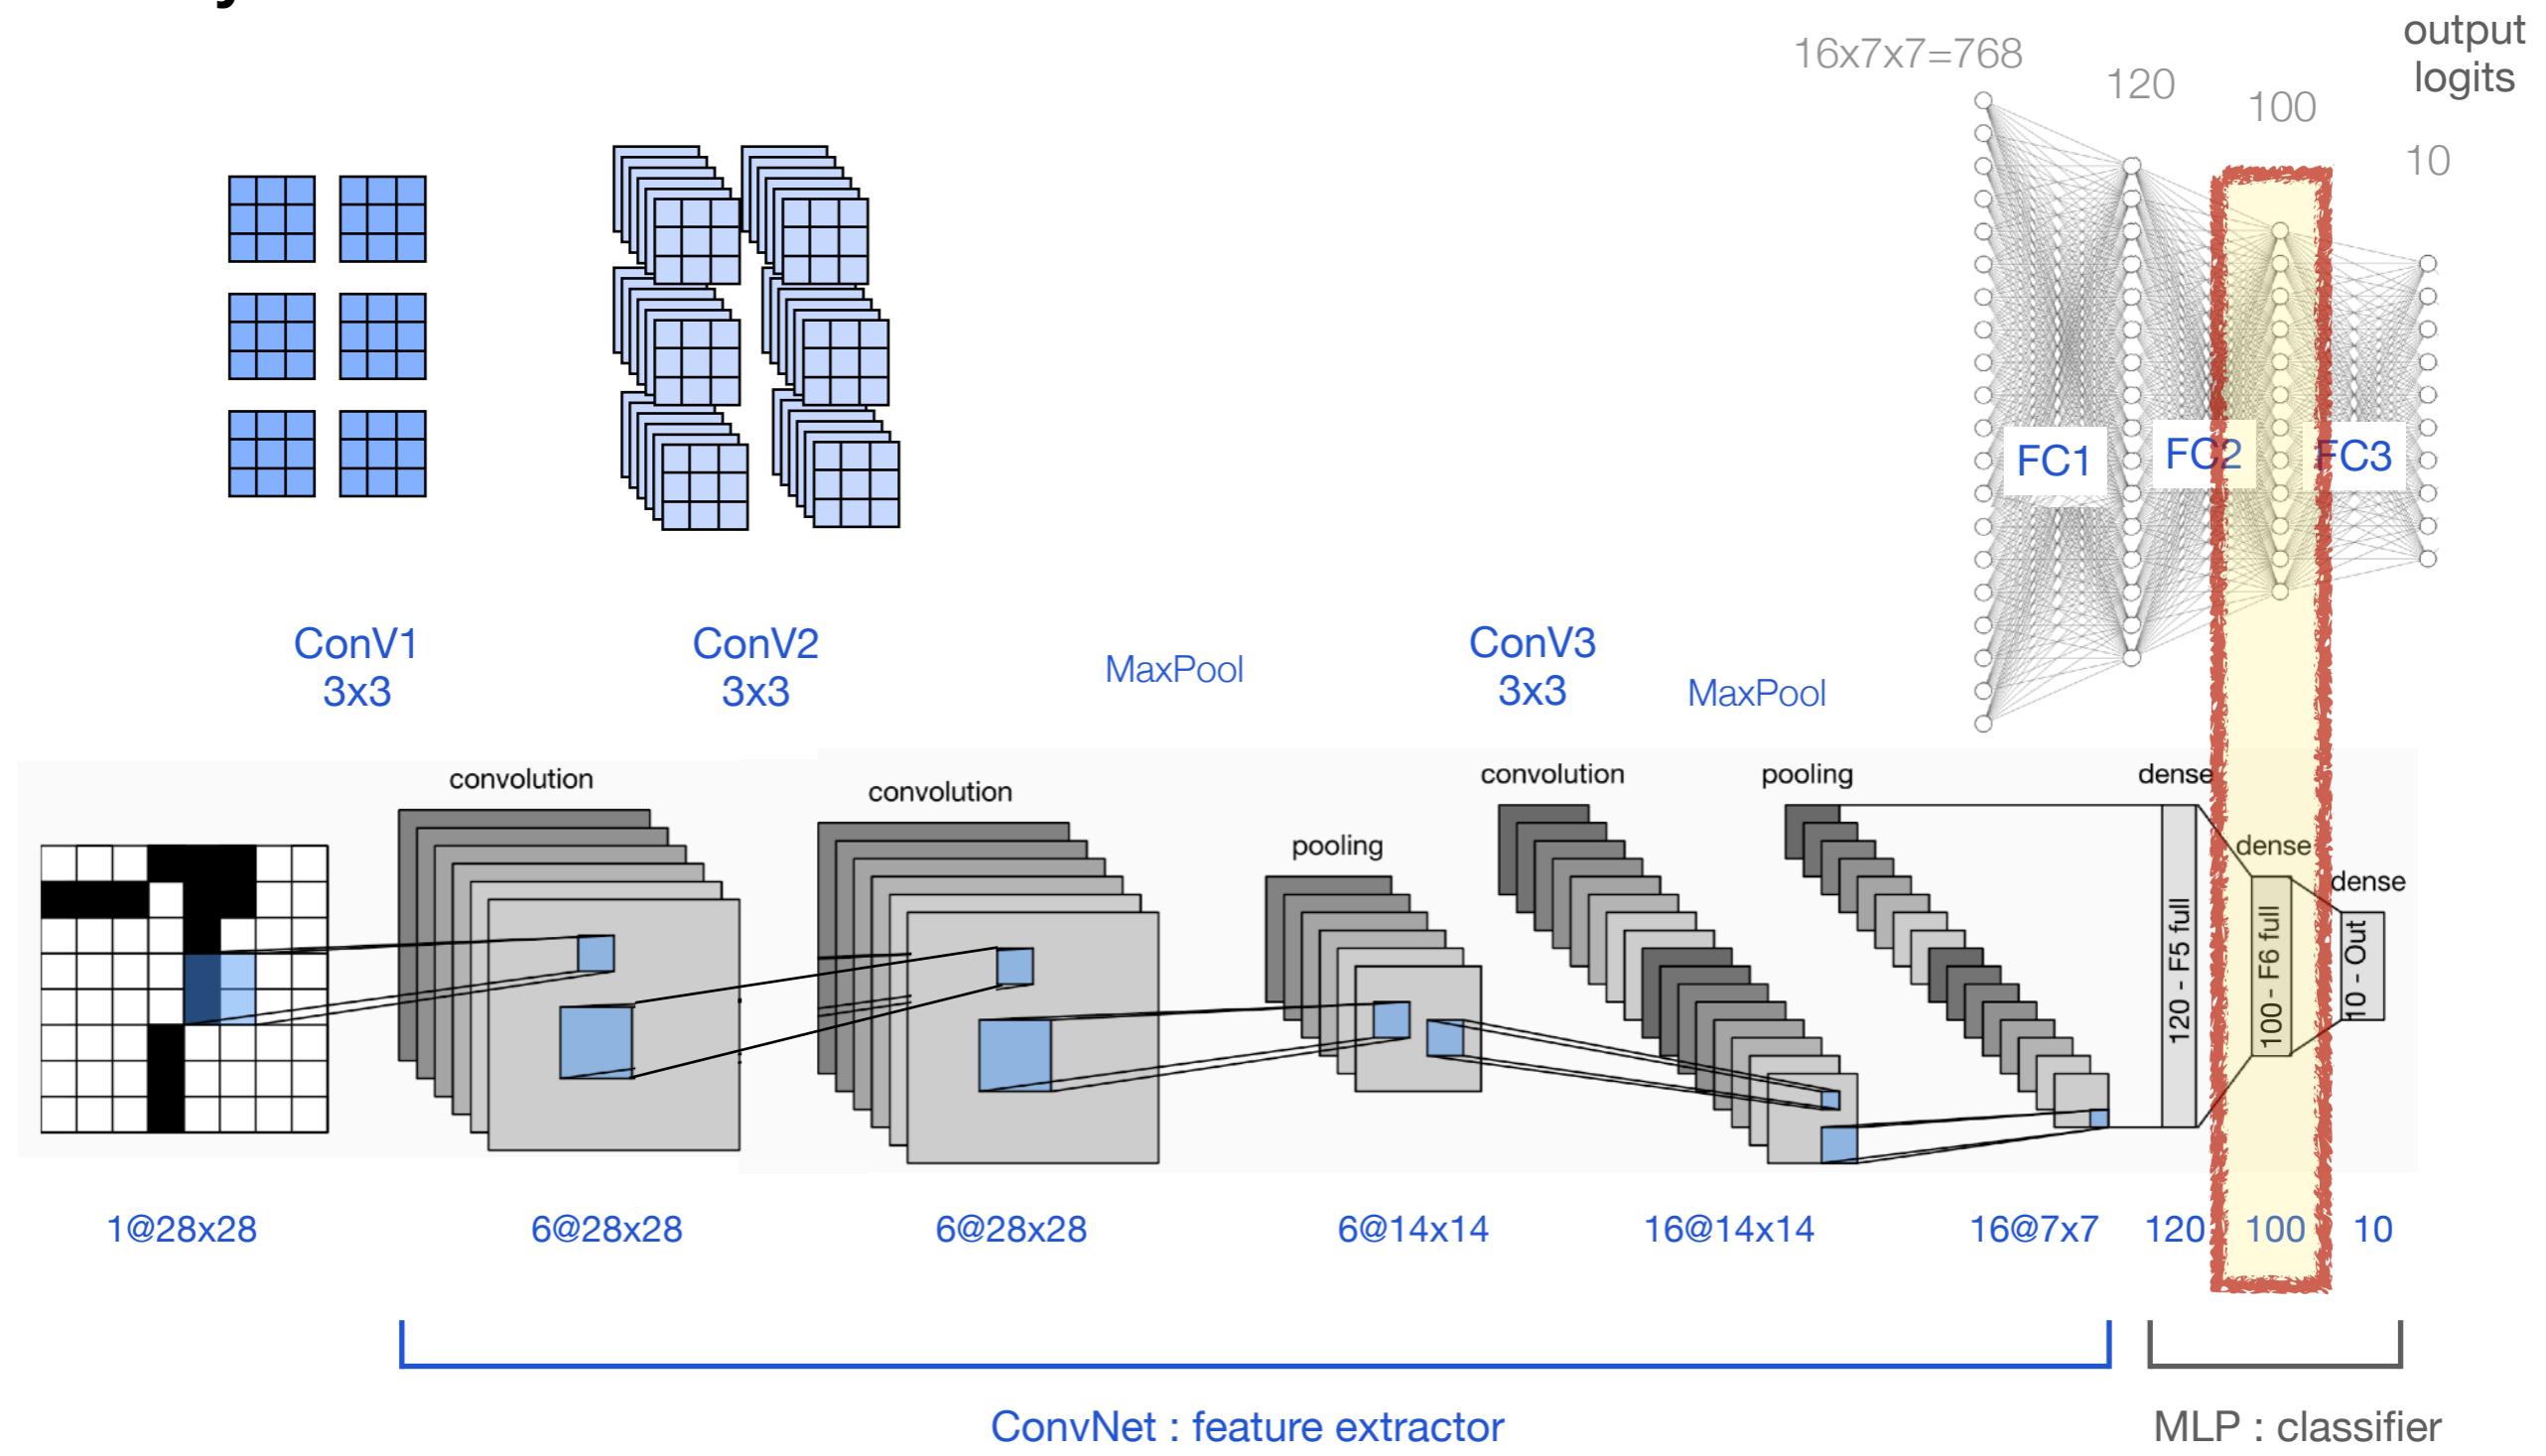
layer 4



layer 5



Last Layer feature vector

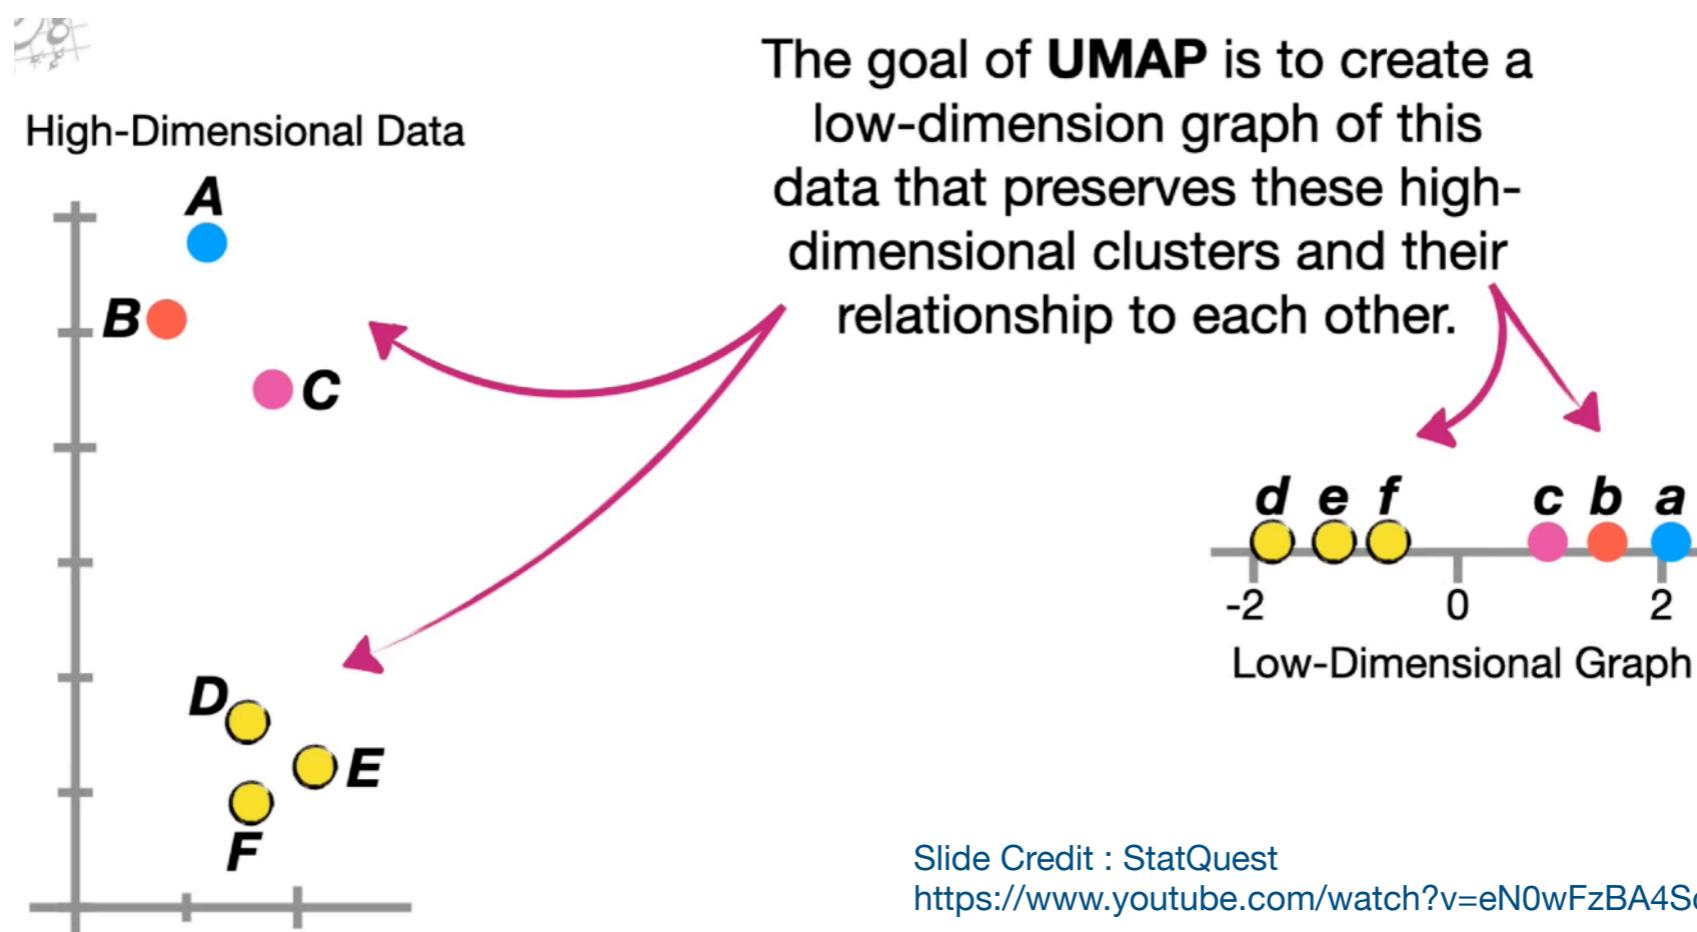


Architecture

ConV1 → ReLU → ConV2 → ReLU → Pool → ConV3 → ReLU → Pool → Flatten → FC1 → ReLU → FC2 → ReLU → FC3 → SoftMax

UMAP – Uniform Manifold Approximation and Projection

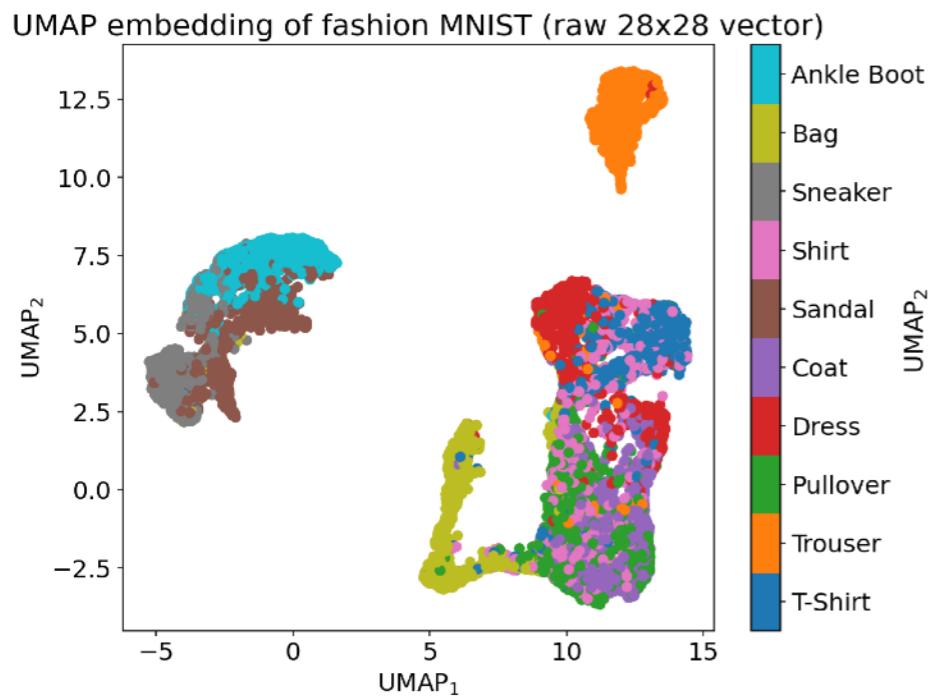
- Our Last Layer Feature Vector has a dimension of 100, for each input image (28x28). How to better visualize such a high dimensional vector ?
- UMAP : an algorithm for dimension reduction that maps high-dimensional vectors to low-dimension while preserving the “similarity” or the “clustering properties” of the original vector.



Clustering of the the Last Layer feature vectors

$$28 \times 28 = 784 \xrightarrow{\text{map}} 2$$

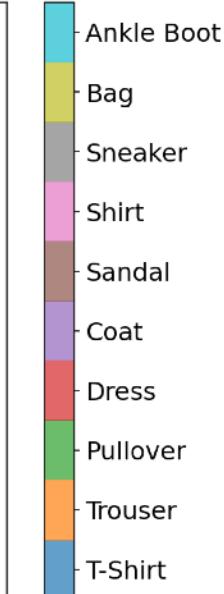
before training



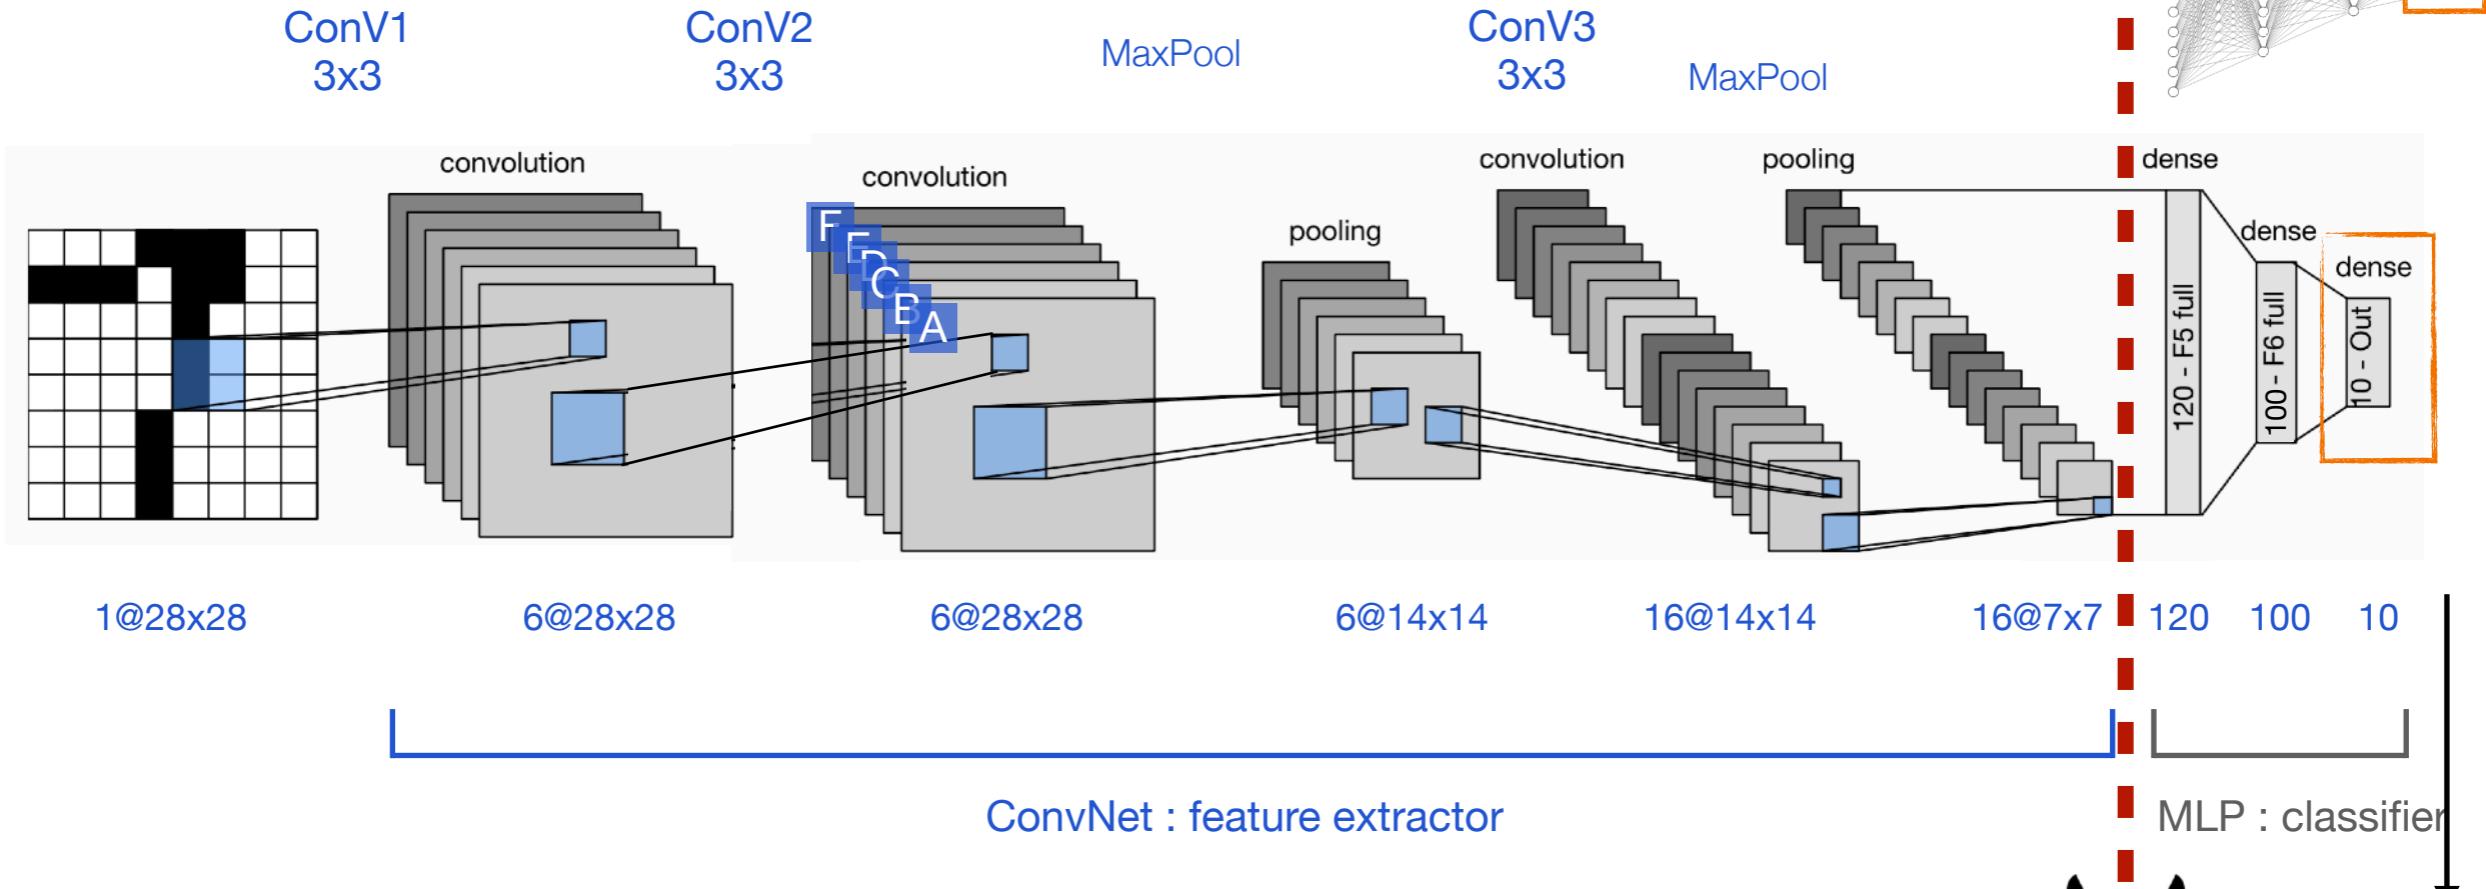
$$100 \xrightarrow{\text{map}} 2$$

epoch=0

epoch=1



Transfer Learning

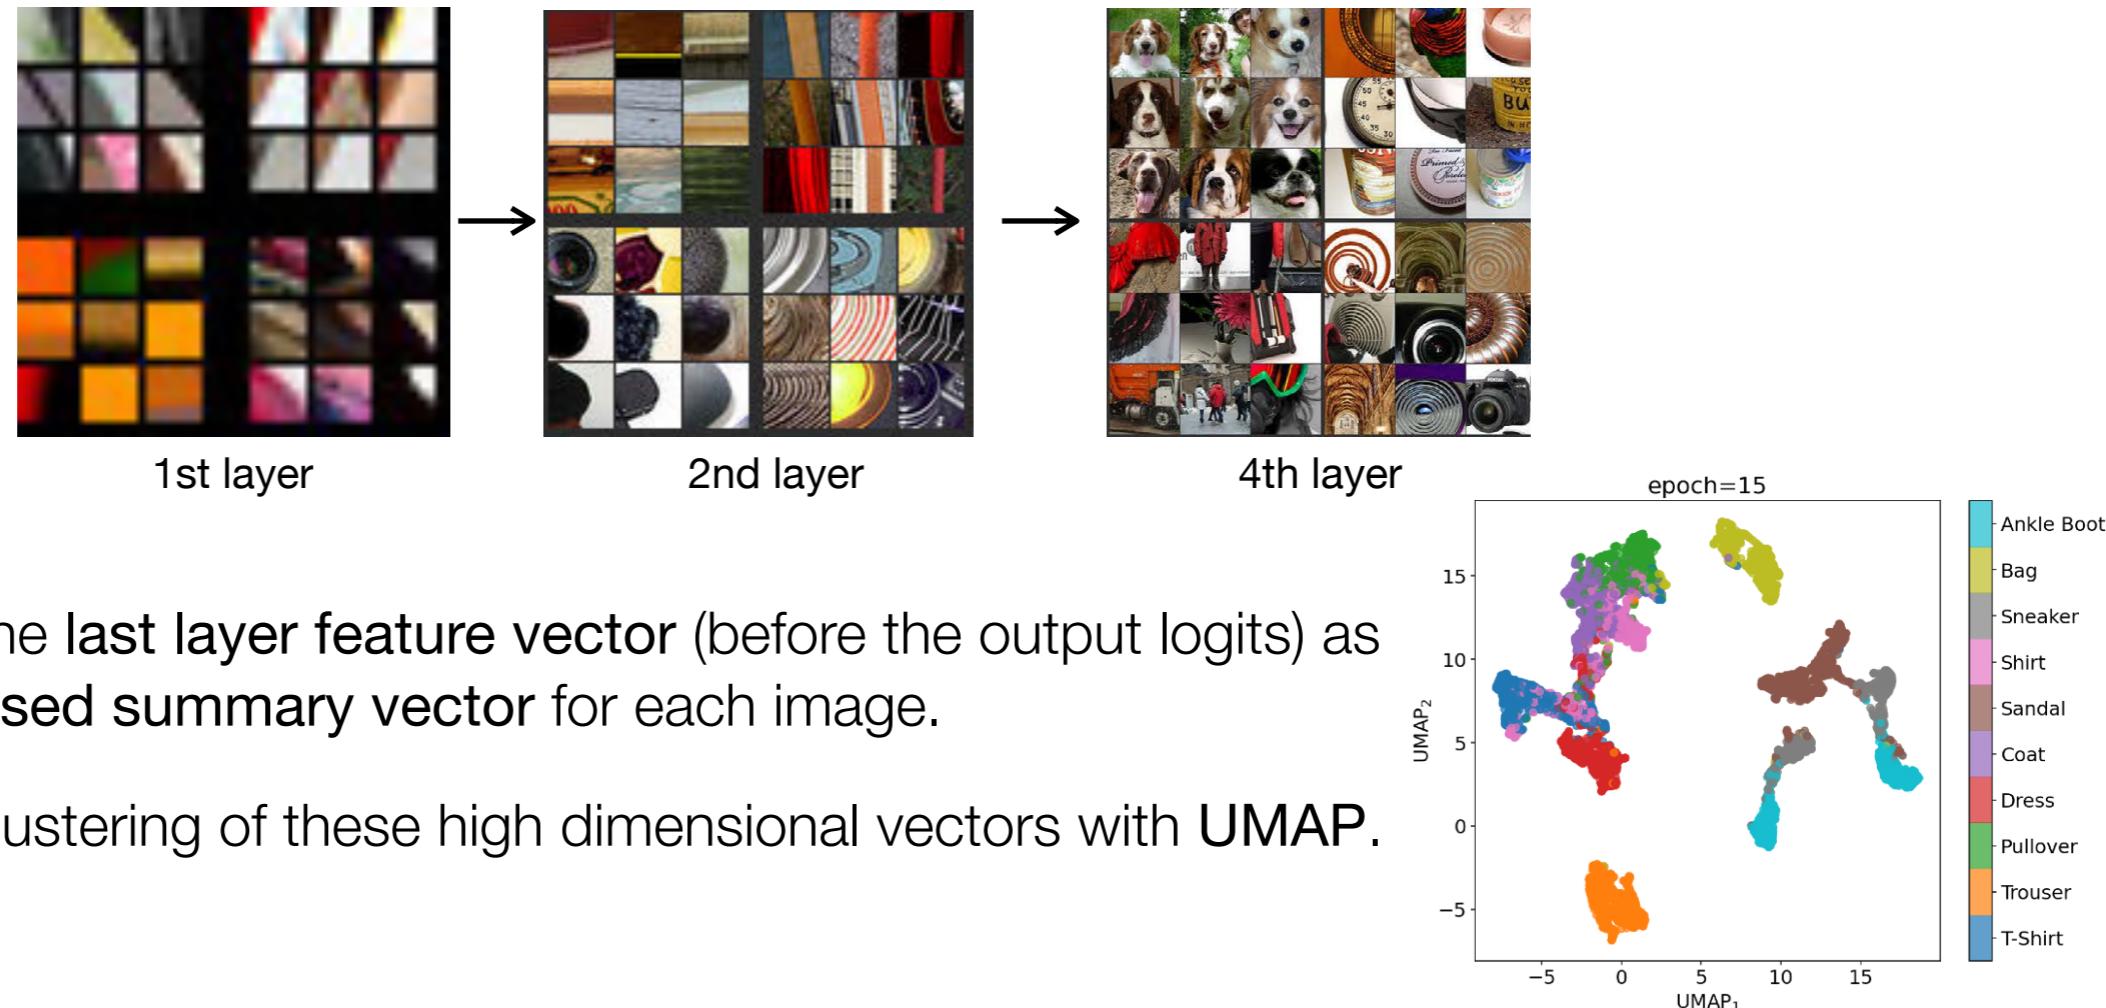


- Take a ConvNet that was trained from scratch with random weight initialization on some large dataset. (e.g. ImageNet with ~1.3 million images)
- Use the convolutional parts of the network as feature extractor.
Freeze the network weight parameters in the convolutional parts.
- Retrain the MLP classification part with the new dataset that you care about.
Modify the MLP network structure if needed (especially the number of output neurons to match the number of outputs for your task).



Summary

- Convolutional Nets learn image patterns hierarchically.
Earlier layers detect simple edges and boundaries.
Composing the lower level features, the later layers learn more complex structures.



- We can view the last layer feature vector (before the output logits) as a final condensed summary vector for each image.

Visualize the clustering of these high dimensional vectors with UMAP.

- Transfer Learning : Utilizing the pre-trained ConvNet as an effective feature extractor.

